

1 **A Bayesian Model of the DNA Barcode Gap**

2 Jarrett D. Phillips^{1,2*} (ORCID: 0000-0001-8390-386X), ... (others?)

3 ¹*School of Computer Science, University of Guelph, Guelph, ON., Canada, N1G2W1*

4 ²*Department of Integrative Biology, University of Guelph, Guelph, ON., Canada, N1G2W1*

5 ***Corresponding Author:** Jarrett D. Phillips¹

6 **Email Address:** jphill01@uoguelph.ca

7 **Running Title:** Bayesian inference for DNA barcode gap estimation

Abstract

A simple statistical model of the DNA barcode gap is outlined. Specifically, accuracy of recently introduced nonparametric metrics, inspired by coalescent theory, to characterize the extent of proportional overlap/separation in maximum and minimum pairwise genetic distances within and among species, respectively, is explored in both frequentist and Bayesian contexts. The empirical cumulative distribution function (ECDF) is utilized to estimate probabilities associated with positively skewed extreme tail distribution quantiles bounded on the closed unit interval $[0, 1]$ based on a straightforward binomial distance overlap count. Using R and the probabilistic programming language, Stan, the proposed maximum likelihood estimators and Bayesian model are demonstrated on CYTB gene sequences from two *Agabus* diving beetle species exhibiting limits in the extent of representative taxonomic sampling. Large-sample theory and MCMC simulations show much uncertainty in parameter estimates, particularly when specimen sample sizes for target species are small. Findings herein highlight the promise of the Bayesian approach using a conjugate beta prior for reliable posterior uncertainty estimation when available data are sparse. Obtained results can shed light on foundational and applied research questions concerning DNA-based specimen identification and species delineation for studies in evolutionary biology and ecology, as well as biodiversity conservation and management, of wide-ranging taxa.

Keywords: Bayesian/frequentist inference, DNA barcoding, intraspecific genetic distance, interspecific genetic distance, specimen identification, species discovery

1 Introduction

The routine use of DNA sequences to support broad evolutionary hypotheses and questions concerning central demographic processes, like gene flow and speciation, is now ubiquitous. Since the late 1980s, mitochondrial DNA (mtDNA) has been employed to produce

a distinctive and measurable pattern of genetic polymorphism in diverse and spatially-distributed taxonomic lineages such as birds, fishes, insects, and arachnids, among other extensively studied groups (Avice et al., 1987). The application of genomic data to applied fields like biodiversity forensics, conservation, and management for the molecular identification of unknown specimen samples came later (*e.g.*, Forensically Informative Nucleotide Sequencing (FINS); Bartlett and Davidson (1992)). Since its inception over 20 years ago, DNA barcoding (Hebert et al., 2003a,b) built significantly on earlier work and has emerged as a robust method of specimen identification and species discovery across myriad multicellular eukaryotes which have been sequenced at easily obtained short, standardized gene regions like the cytochrome *c* oxidase subunit I (5'-COI) mitochondrial locus for animals. However, DNA barcoding is not without its controversies.

The success of the single-locus approach, particularly for regulatory and forensic applications, depends crucially on two important factors: (1) the availability of high-quality specimen records found in public reference sequence databases such as the Barcode of Life Data Systems (BOLD; <http://www.barcodinglife.org>) (Ratnasingham and Hebert, 2007) and GenBank (<https://www.ncbi.nlm.nih.gov/genbank/>), and (2) the establishment of a DNA barcode gap — the notion that the maximum genetic distance observed within species is much smaller than the minimum degree of marker variation found among species (Meyer and Paulay, 2005; Meier et al., 2008). Early work has demonstrated that the presence of a DNA barcode gap hinges strongly on extant levels of species haplotype diversity gauged from comprehensive specimen sampling at wide geographic and ecological scales (Bergsten et al., 2012; Čandek and Kuntner, 2015). Despite this, many taxonomic groups lack adequate separation in their pairwise intraspecific and interspecific genetic distances due to varying rates of evolution in both genes and taxa (Pentinsaari et al., 2016). Furthermore, it has been well-demonstrated that the presence of a DNA barcode gap becomes less certain with increasing spatial scale of sampling, since interspecific distances increase due to population structuring under limited dispersal through isolation-by-distance, while intraspecific distances

shrink as more closely-related species are sampled (Phillips et al., 2022). This can pose problems in cases of rare singleton species or monotypic taxa, for instance (Ahrens et al., 2016). As a result, rapid matching of unknown samples to expertly-validated references can be compromised, leading to cases of false positives (taxon oversplitting) and false negatives (excessive lumping of taxa) due to incomplete lineage sorting, interspecies hybridization, genome introgression, species synonymy, cryptic species diversity, and misidentifications (Hubert and Hanner, 2015; Phillips et al., 2022). Another culprit includes PCR/sequencing error caused by incorrect nucleotide basecalling, chimeras/heteroplasmies, sample contamination, indels (insertion/deletions), NUMT (nuclear-mitochondrial insert) pseudogene amplification, or *Wolbachia* endosymbiont infection, among other explanations (Phillips et al., 2023). Each of these can artificially magnify levels of genetic diversity, rendering techniques like DNA barcoding unreliable for certain groups like marine invertebrates, which are both undersampled and possess slow rates of molecular evolution (Bucklin et al., 2011).

Recent work has argued that DNA barcoding, in its current form, is lacking in statistical rigor (Phillips et al., 2022). Most studies rely strongly on heuristic distance-based thresholds calculated from convenient specimen sample sizes, rather than on optimal sampling designs using collated taxon haplotype information (Phillips et al., 2015, 2020, 2022), to infer taxonomic identity and delineate species boundaries on the basis of clustering variation into putative groups reminiscent of species. Of these studies, few report measures of uncertainty, such as standard errors (SEs) and confidence intervals (CIs), around estimates of intraspecific and interspecific variation (*e.g.*, Wiemers and Fiedler (2007)), calling into question the existence of a true species' DNA barcode gap (Čandek and Kuntner, 2015; Phillips et al., 2022). To support this idea, novel nonparametric locus- and species-specific metrics based on the multispecies coalescent (MSC) were recently outlined by Phillips et al. (2024). The basic coalescent (Kingman, 1982a,b) encompasses a backwards continuous-time stochastic Markov process of allelic sampling among lineages under the infinite sites model of mutation

88 within a large neutrally-evolving, panmictic, diploid species population of constant size from
 89 the present into the past towards the most recent common ancestor (MRCA). Although the
 90 coalescent sampling process plays a foundational role in evolutionary biology and population
 91 genetics theory due to its inherent simplicity and flexibility, the approach has been both
 92 under-utilized and under-appreciated as a key player within DNA barcoding (Collins and
 93 Cruickshank, 2014; Downton et al., 2014; Hubert and Hanner, 2015; Phillips et al., 2022).
 94 Numerous theoretical and empirical investigations have shown that protein-coding sequence
 95 variation within haploid loci like COI is generally subject to strong purifying selection and
 96 selective sweeps at nonsynonymous sites, rather than the occurrence of
 97 functionally silent substitutions within synonymous regions (Stoeckle and Thaler, 2014). In
 98 effect, accumulated mutations become diagnostic of species at a much faster rate compared
 99 to nuclear regions of the genome (Hubert and Hanner, 2015; Phillips et al., 2022). Previously
 100 proposed MSC algorithmic approaches (of which there are too many to exhaustively list here,
 101 and which have mostly been applied to multilocus sequence data), generally assume a strict
 102 molecular clock and a simplified model of DNA sequence evolution across closely-related taxa,
 103 from which an estimated species phylogeny may be constructed (*e.g.*, with or without use of
 104 a guide tree) (*e.g.*, Rannala and Yang (2003, 2017); Yang and Rannala (2010, 2014, 2017)).
 105 Even tree- and network-based taxon delimitation methods designed for single loci, such as
 106 the Generalized Mixed Yule Coalescent (GMYC) (Pons et al., 2006; Monaghan et al., 2009;
 107 Fujisawa and Barraclough, 2013), Poisson Tree Processes (PTP) (Zhang et al., 2013; Kapli
 108 et al., 2017), and statistical parsimony/haplotype networks (*e.g.*, TCS networks (Templeton
 109 et al., 1992)), which have seen much use in DNA barcoding initiatives, have their flaws.
 110 Optimal performance of these methods relies heavily on careful parameter selection (*e.g.*,
 111 prior specification of branching rates in ultrametric trees in the case of GMYC) (Hubert
 112 et al., 2024) within third party software like BEAST (Bouckaert et al., 2019), among other
 113 concerns (Talavera et al., 2013; Fonseca et al., 2021). Careless parameterization of established
 114 approaches by users unfamiliar with their underlying statistical characteristics may severely

bias overall results in unintended ways (Hart and Sunday, 2007). In contrast, Phillips et al.'s (2024) approach is both tree- and network-free and does not require judicious parameter setting. Therefore, the method is extremely efficient and fast to run.

The DNA barcode gap statistics have been shown to hold strong promise for reliable DNA barcode gap assessment when applied to predatory *Agabus* (Coleoptera: Dytiscidae) diving beetles (Phillips et al., 2024). With a Holarctic range, despite their ease of sampling and well-established taxonomy, this group possesses few morphologically-distinct taxonomic characters that readily facilitate their assignment to the species level (Bergsten et al., 2012). Accurate taxon discrimination within *Agabus* is paramount since several species are currently classified as endangered on the International Union for Conservation of Nature (IUCN) Red List (<https://www.iucnredlist.org>). Further, the proposed metrics indicate that sister species pairs from this taxon are often difficult to distinguish on the basis of their DNA barcode sequences (Phillips et al., 2024), which may be caused by *Agabus*'s increased degree of non-monophyly at wider geospatial scales, occurrence of hybridization/introgression, or presence of synonymy (Bergsten et al., 2012). Using sequence data from three mitochondrial cytochrome markers (5'-COI, 3'-COI, and cytochrome *b* (CYTB)) obtained from BOLD and GenBank, results from Phillips et al. (2024) highlight that DNA barcoding has been a one-sided argument. Phillips et al.'s (2024) findings point to the need to balance both the sufficient collection of specimens, as well as the extensive sampling of species. Not surprisingly, DNA barcode libraries for most taxonomic groups are biased toward the latter effort (Phillips et al., 2015, 2019, 2022).

The estimators from Phillips et al. (2024) represent a clear improvement over simple, yet arbitrary, distance heuristics such as the 2% rule noted by Hebert et al. (2003a) and the 10 $\times$ rule (Hebert et al., 2004) that form the core of threshold-based single-locus species delimitation tools like Automatic Barcode Gap Discovery (ABGD) (Puillandre et al., 2011), Assemble Species by Automatic Partitioning (ASAP) (Puillandre et al., 2021), and the Barcode Index Number (BIN) framework (Ratnasingham and Hebert, 2013). The 2% rule

142 asserts that DNA barcodes differing by at least 2% at sequenced genomic regions should be
 143 expected to originate from different biological species, whereas the 10 $\times$ rule suggests that
 144 barcode sequences displaying 10 times more genetic variation among species than within taxa
 145 is suggestive of a distinct evolutionary origin. However, the lack of adoption of an explicit
 146 and universally agreed upon species concept in DNA barcoding studies that readily governs
 147 lineage formation and proliferation necessary to establish rigorous taxon definitions for
 148 successful delimitation of hypothesized and heuristic evolutionary units using these
 149 well-known criteria, in conjunction with secondary lines of evidence (*e.g.*, morphology,
 150 ecology, geography, ontogeny, and behaviour) promised by an integrative framework (Dayrat,
 151 2005), is missing (de Queiroz, 2007; Rannala, 2015; Pante et al., 2015; Wells et al., 2022). This
 152 is the case, for instance, with species existing in allopatry/sympatry, where establishment of
 153 reproductive isolation on the basis of fixed diagnostic differences using the Biological Species
 154 Concept (BSC) (Mayr, 1942), is challenging (Fujita et al., 2012). In addition, the reliance on
 155 visualization approaches, such as frequency histograms, dotplots, and quadrant plots, to
 156 expose DNA barcoding’s limitations, has also been criticized since they can fog real species
 157 signal (Collins and Cruickshank, 2013; Phillips et al., 2022). Up until the work of Phillips
 158 et al. (2024), the majority of studies (*e.g.*, Young et al. (2021)) have treated the DNA barcode
 159 gap as a binary response. Given poor sampling depth for most taxa, a Yes/No dichotomy
 160 is inherently flawed because it can falsely imply a DNA barcode gap is present for a taxon
 161 of interest when in fact no such separation in genetic distances exists. On the other hand,
 162 in the case of sequence overlap, while the presence a single specimen record is a necessary
 163 requirement to nullify the existence of a DNA barcode gap in light of available data, it is not a
 164 sufficient condition, since said taxon record could represent a species misidentification or other
 165 error. Thus, the MSC represents an ideal model to settle ongoing debates and issues regarding
 166 DNA barcoding’s use as both a molecular identification and delineation tool. Notably, Fujita
 167 et al. (2012) remarks that the MSC applied to the species delimitation problem harmonizes
 168 many species concepts since it identifies independent evolutionary trajectories under the

umbrella of the General Lineage Concept. The recently introduced DNA barcode gap metrics go some way into operationalizing this task.

Phillips et al.'s (2024) proposed statistics quantify the extent of asymmetric directionality of proportional distance distribution overlap/separation for species within well-sampled taxonomic genera based on a straightforward distance count, in a similar vein to established measures of statistical similarity such as the Kullback-Leibler (KL) divergence (Kullback and Leibler, 1951) and other related statistics. The metrics can be employed in a variety of ways, including to validate performance of marker genes for specimen identification to the species level (using nonparametric bootstrapping, as in Phillips et al. (2024)). The estimators can further be utilized to assess whether computed values are consistent with fundamental population genetic-level parameters like effective population size (N_e), mutation rates (μ), migration rates (m), and divergence times (τ), as well as derived ones, such as the population mutation rate $\theta = 4N_e\mu$ and the fixation index, F_{ST} , for species under study in a statistical phylogeographic setting (Knowles and Maddison, 2002; Knowles, 2009; Mather et al., 2019). Early on, species discovery using DNA barcoding was presumed to only work for reciprocally monophyletic groups, and thus concerned itself with terminal branches of generated phylogenies rather than more basal lineages occurring deeper in hypothesized species trees (Fujita et al., 2012; Mutanen et al., 2016). Furthermore, the occurrence of short branches and large N_e within resolved phylogenies increases the probability of deep coalescence, clouding species delimitations, which often fail or are uncertain in broad parameter space (Carstens et al., 2013; Hickerson et al., 2006; Rannala, 2015). As DNA barcoding is a single-locus approach, it is problematic for evolutionarily young taxa, wherein incomplete lineage sorting within gene genealogies is a common phenomenon due to the ongoing stochastic dynamic of mutation generating population variation, and genetic drift driving gene variants to fixation (Rannala, 2015). The most promising way forward in assessing the validity of Phillips et al.'s (2024) method seems to be through the use of software such as BPP (Bayesian Phylogenetics and Phylogeography), which permits efficient

full Bayesian simulations under various MSC models (*e.g.*, MSC-I (MSC with introgression) or MSC-M (MSC with migration), among others) using MCMC for tree parameter estimation (via the A00 option, for instance) (Flouri et al., 2018), or PHRAPL (Phylogeographic Inference using Approximate Likelihoods) (Jackson et al., 2017a,b), which employs tractable phylogenetic likelihood calculations via the genealogical divergence index (gdi). Other hypothesis-driven approaches to species delimitation model selection, like Bayes Factors (BFs) (Leaché et al., 2014), are also possible.

While introduction of the new metrics is a step in the right direction, what appears to be missing is a rigorous statistical treatment of the DNA barcode gap. This includes an unbiased way to compute the statistical accuracy of Phillips et al.’s (2024) estimators arising through problems inherent in frequentist maximum likelihood estimation for probability distributions having bounded positive support on the closed unit interval $[0, 1]$. To this end, here, a Bayesian model of the DNA barcode gap coalescent is introduced to rectify such issues. The model allows accurate estimation of posterior means, posterior standard deviations (SDs), posterior quantiles, and credible intervals (CrIs) for the statistics given datasets of intraspecific and interspecific distances for species of interest.

2 Methods

2.1 DNA Barcode Gap Metrics

The novel nonparametric maximum likelihood estimators (MLEs) of proportional overlap/separation between intraspecific and interspecific distance distributions for a given species (x) to aid assessment of the DNA barcode gap are as follows:

$$p_x = \frac{\#\{d_{ij} \geq a\}}{\#\{d_{ij}\}} \quad (1)$$

$$q_x = \frac{\#\{d_{XY} \leq b\}}{\#\{d_{XY}\}} \quad (2)$$

$$p'_x = \frac{\#\{d_{ij} \geq a'\}}{\#\{d_{ij}\}} \quad (3)$$

$$q'_x = \frac{\#\{d'_{XY} \leq b\}}{\#\{d'_{XY}\}} \quad (4)$$

where d_{ij} are distances within species, d_{XY} are distances among species for an *entire* genus of concern, and d'_{XY} are combined interspecific distances for a target species and its closest neighbouring species. The notation $\#$ reflects a count. Quantities a , a' , and b correspond to $\min(d_{XY})$, $\min(d'_{XY})$, and $\max(d_{ij})$, the minimum interspecific distance, the minimum combined interspecific distance, and the maximum intraspecific distance, respectively (Figure 1).




Figure 1: Modified depiction from Meyer and Paulay (2005) and Phillips et al. (2024) of the overlap/separation of intraspecific and interspecific distances (δ) for calculation of the DNA barcode gap metrics (p_x and q_x) for a hypothetical species x . The minimum interspecific distance is denoted by a and the maximum intraspecific distance is indicated by b . The quantity $f(\delta)$ is akin to a kernel density estimate of the probability density function of distances. A similar visualization can be displayed for p'_x and q'_x within the interval $[a', b]$.

Hence, Equations (1)-(4) are simply empirical partial means of distances falling at and below,

or at and exceeding, given distribution thresholds. Notice further that a/a' , and b are also
 the first and n th order statistics, $X_{(1)}$ and $X_{(n)}$, respectively, with $a/a' < b$, which have been
 pointed out by Phillips et al. (2022) as important for developing a statistical theory to test for
 the existence of the DNA barcode gap. Equations (1)-(4) can also be expressed in terms of
 empirical cumulative distribution functions (ECDFs) (see next section). A clear connection
 between the ECDF and the nonparametric MLE was noted by Efron and Tibshirani (1993)
 as well as others. Distances form a continuous distribution and are easily computed from a
 model of DNA sequence evolution, such as uncorrected or corrected p-distances (Jukes and
 Cantor, 1969; Kimura, 1980) using, for example, the `dist.dna()` function available in the
ape R package (Paradis et al., 2004). However, calculated values are *not* independent and
 identically distributed (IID) because estimated standard errors (SEs) will depend on both the
 number of species sampled within the genus under study, as well as the number of specimens
 sampled within a target species. In Phillips et al. (2024), To tease this out, Phillips et al.
 (2024) suggests plotting estimator values against their estimated SEs, along with a simple
 random downsampling scheme. In the case of two species comprising a focal genus, one well
 sampled and the other poorly sampled, values of the metrics close to zero for the sufficiently
 sampled species will likely possess larger SEs following downsizing to match the number of
 poorly sampled specimens (Phillips et al., 2024). The approach of Phillips et al. (2024) differs
 markedly from the traditional definition of the DNA barcoding gap laid out by Meyer and
 Paulay (2005) and Meier et al. (2008) in that the proposed metrics incorporate interspecific
 distances which *include* the target species of interest. Furthermore, if a focal species is
 found to have multiple nearest neighbours, then the species possessing the smallest average
 distance is used (Phillips et al., 2024). These schemes more accurately account for species'
 coalescence histories inferred from contemporaneous samples of DNA sequences leading to
 instances of barcode sequence sharing, such as interspecific hybridization/introgression events
 (Phillips et al., 2024). For example, within Equations (3) and (4), the degree of distance
 distribution overlap/separation between a target taxon and its nearest neighbouring species,

gauged from magnitudes of p'_x and q'_x , is directly proportional to the amount of time in which the two lineages diverged from the MRCA (Phillips et al., 2024). More specifically, $a \propto t_{XY, \min}$ and $b \propto t_{ij, \max}$, the minimum and maximum (unobserved) divergence time (in N_e generations), respectively, for two different individuals (i, j) and species (X, Y) within a well-inventoried genus (Phillips et al., 2024). Figure 1 in Phillips et al. (2024) depicts a phylogeny for two hypothetical species, A and B , showing paraphyly based on the sequencing of two mitochondrial loci. The (unobserved) time the two species formed, t_{AB} , is intermediate between $t_{XY, \min}$ and $t_{ij, \max}$, but theoretically, its maximum ($t_{AB, \max}$) may occur more distantly in the past closer to the root comprising the MRCA for all sampled taxa (Phillips et al., 2024). From a coalescence perspective, p_x is the probability two lineages within species x coalesce after $t_{XY, \min}$, whereas q_x represents the probability a lineage within species x coalesces with a lineage in species k (within a sample) after $t_{XY, \min}$, but before $t_{ij, \max}$. Thus, the quantities can be used as a criterion to assess the failure of DNA barcoding in recently radiated taxonomic groups, among other plausible biological explanations mentioned earlier. Note, distances are constrained to the interval $[0, 1]$, whereas the metrics are defined only on the interval $[a/a', b]$. Values of the estimators obtained from Equations (1)-(4) close to or equal to zero give evidence for separation between intraspecific and interspecific distance distributions; that is, values suggest the presence of a DNA barcode gap for a target species. Conversely, values near or equal to one give evidence for distribution overlap; that is, values likely indicate the absence of a DNA barcode gap.

2.2 The Model

Before delving into the derivation of the proposed DNA barcode gap metrics, review of some fundamental statistical theory is necessary.

For a given random variable X from a given distribution F , its cumulative distribution function (CDF) is defined by

$$F_X(t) = \mathbb{P}(X \leq t) = 1 - \mathbb{P}(X > t). \quad (5)$$

276 Rearranging Equation (5) gives

$$\mathbb{P}(X > t) = 1 - F_X(t) = 1 - \mathbb{P}(X \leq t), \quad (6)$$

277 from which it follows that

$$\mathbb{P}(X \geq t) = 1 - F_X(t) + \mathbb{P}(X = t). \quad (7)$$

278 The CDF has many desirable statistical properties, such as boundedness, continuity, and
 279 monotonicity. Consult any graduate level text on mathematical statistics for reference.

280 Equations (1)-(4) can thus be written in terms of ECDFs via Equations (5) and (7) as
 281 follows, since the true underlying F s, are unknown *a priori*, and therefore must be estimated
 282 using available sequence distance data:

$$\begin{aligned}
p_x &= \mathbb{P}(d_{ij} \geq a) \\
&= 1 - \hat{F}_{d_{ij}}(a) + \mathbb{P}(d_{ij} = a)
\end{aligned} \tag{8}$$

$$\begin{aligned}
q_x &= \mathbb{P}(d_{XY} \leq b) \\
&= \hat{F}_{d_{XY}}(b)
\end{aligned} \tag{9}$$

$$\begin{aligned}
p'_x &= \mathbb{P}(d_{ij} \geq a') \\
&= 1 - \hat{F}_{d_{ij}}(a') + \mathbb{P}(d_{ij} = a')
\end{aligned} \tag{10}$$

$$\begin{aligned}
q'_x &= \mathbb{P}(d'_{XY} \leq b) \\
&= \hat{F}_{d'_{XY}}(b).
\end{aligned} \tag{11}$$

283 Given n increasing-ordered data points, the (discrete) ECDF, $\hat{F}_n(t) = \frac{1}{n} \sum_{i=1}^n \mathbb{1}_{[x_i \leq t]}$,
 284 comprises an increasing step function having jump discontinuities of size $\frac{1}{n}$ at each sample
 285 observation (x_i) , excluding ties (or steps of weight $\frac{i}{n}$ with duplicate observations), where $\mathbb{1}(x)$
 286 is the indicator function. Note that $1 - \hat{F}_n(t)$ is the complementary ECDF (CECDF), which
 287 is a declining function. Further, notice that $\mathbb{P}(X = t) \neq 0$. The ECDF/CECDF has good
 288 asymptotic behaviour in the limit of sample size. Equations (8)-(11) clearly demonstrate
 289 the asymmetric directionality of the proposed metrics. Plots of the ECDFs/CECDFs for
 290 nonzero intraspecific, interspecific, and combined interspecific distances, superimposed with
 291 corresponding values of a/a' and b , provide a new, useful way to visualize the DNA barcode
 292 gap. Firstly, since there is no need to tune the output (*e.g.*, number of bins and bin width
 293 in the case of histograms). Secondly, probabilities reflecting areas under the distance curves
 294 depicted in **Figure 1** can be directly read off the ECDF/CECDF curves. Values of the metrics
 295 can also be determined from direct inspection of the ECDF/CECDF plot by determining the
 296 corresponding y -value (cumulative probability) for a given value of x (distance). Furthermore,
 297 calculation of the DNA barcode gap estimators is convenient as they implicitly account for

total distribution area (including overlap).

A major criticism of large sample (frequentist) theory is that it relies on asymptotic properties of the MLE (whose population parameter is assumed to be a fixed but unknown quantity), such as estimator normality and consistency as the sample size approaches infinity. This problem is especially pronounced in the case of binomial proportions (Newcombe, 1998). The estimated Wald SE of the sample proportion, is given by $\widehat{SE}[\hat{p}] = \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$, where $\hat{p} = \frac{Y}{n}$ is the MLE, Y is the total number of successes ($Y = \sum_{i=1}^n y_i$), and n is the total number of trials (*i.e.*, sample size). However, the above formula for the standard error is problematic for several reasons. First, it is a Normal approximation which makes use of the central limit theorem (CLT); thus, large sample sizes are required for reliable estimation. When few observations are available, SEs will be large and inaccurate, leading to low statistical power to detect a true DNA barcode gap when one actually exists. Further, resulting interval estimates could span values less than zero or greater than one, or have zero width, which is practically meaningless. Second, when proportions are exactly equal to zero or one, resulting SEs will be exactly zero, rendering $\widehat{SE}[\hat{p}]$ given above completely useless. In the context of the proposed DNA barcode gap metrics, values obtained at the boundaries of their support are often encountered. Therefore, reliable calculation of SEs is not feasible. Given the importance of sufficient sampling of species genetic diversity for DNA barcoding initiatives, a different statistical estimation approach is necessary.

Bayesian inference offers a natural path forward in this regard since it allows for straightforward specification of prior beliefs concerning unknown model parameters and permits the seamless propagation of uncertainty, when data are lacking and sample sizes are small, through integration with the likelihood function associated with true generating processes. The posterior distribution ($\pi(\theta|Y)$) is given by Bayes' theorem up to a proportionality $\pi(\theta|Y) \propto \pi(Y|\theta)\pi(\theta)$, where θ are unobserved parameters, Y are known data, $\pi(Y|\theta)$ is the likelihood, and $\pi(\theta)$ is the prior. As a consequence, because parameters are treated as random variables, Bayesian models are much more flexible and generally more

easily interpretable compared to frequentist approaches. Under the Bayesian paradigm, entire posterior distributions, along with their summaries (*e.g.*, CrIs) are outputted, rather than just long run behaviour reflected in sampling distributions, p-values, and CIs as in the frequentist case, thus allowing direct probability statements to be made.

Essentially, from a statistical perspective, the goal herein is to nonparametrically estimate probabilities corresponding to extreme tail quantiles for positive highly skewed distributions on the unit interval (or any closed subinterval thereof). Here, it is sought to numerically approximate the extent of proportional overlap/separation of intraspecific and interspecific distance distributions within the subinterval $[a/a', b]$. This is a challenging computational problem within the current study as detailed in subsequent sections. The usual approach employs kernel density estimation (KDE), along with numerical or Monte Carlo integration of complicated probability distribution functions (PDFs), and invocation of extreme value theory (EVT); however, this requires careful selection of the bandwidth parameter, among other considerations. This becomes problematic when fitting finite mixture models where nonidentifiability is rampant. For DNA barcode gap estimation, this would correspond to a two-component mixture (one for intraspecific distance comparisons, and the other for interspecific comparisons), with one or more distribution modes or curve intersection points between components, and the presence of zero distance inflation. This makes parameter estimation difficult using methods like the Expectation-Maximization (EM) algorithm (Dempster et al., 1977) as the algorithm may become stuck in suboptimal regions of the parameter search space and prematurely converge to local optima. This seems wholly plausible, especially given that the genus-level distribution of genetic distances is a mixture of closely-related (*i.e.*, nearest neighbour) and more distant species. Here, for simplicity, an alternate route is taken to avoid these obstacles.

Counts, y , of overlapping distances (as expressed in the numerator of Equations (1)-(4)) are treated as binomially distributed with expectation $\mathbb{E}[Y|\theta] = k\theta$, and variance $\mathbb{V}[Y|\theta] = k\theta(1 - \theta)$, where $k = \{N, C\}$ are total count vectors of intraspecific and combined

interspecific distances, respectively, for a target species along with its nearest neighbour species, and $k = M$ is a total count vector for all interspecific species comparisons. This follows from the fact that the ECDF/CECDF is binomially distributed. The quantity thus being estimated is the parameter vector $\theta = \{p_x, q_x, p'_x, q'_x\}$. The metrics encompassing θ are presumed to follow a Beta(α, β) distribution, with real shape parameters α and β , which is a natural choice of prior on probabilities. The beta distribution has a prior mean of $\mathbb{E}[\theta|\alpha, \beta] = \frac{\alpha}{\alpha+\beta}$ and a prior variance equal to $\mathbb{V}[\theta|\alpha, \beta] = \frac{\alpha\beta}{(\alpha+\beta)^2(\alpha+\beta+1)}$. In the case where $\alpha = \beta$, all generated Beta(α, β) distributions will possess the same prior expectation, whereas the prior variance will shrink as both α and β increase. Such a scheme is quite convenient since the beta distribution is conjugate to the binomial distribution. Thus, the posterior distribution is also beta distributed, specifically, Beta($\alpha + Y, \beta + n - Y$), having expectation $\mathbb{E}[\theta|Y, \alpha, \beta] = \frac{\alpha+Y}{\alpha+\beta+n}$ and variance $\mathbb{V}[\theta|Y, \alpha, \beta] = \frac{(\alpha+Y)(\beta+n-Y)}{(\alpha+\beta+n)^2(\alpha+\beta+n+1)}$. In the context of DNA barcoding, it is important that the DNA barcode gap metrics effectively differentiate between extremes of no overlap/complete separation and complete overlap/no separation, corresponding to values of the metrics equal to 0 and 1 (equivalent to total distance counts of 0 and n), respectively. These extremes yield a posterior expectation of $\mathbb{E}[\theta|Y = 0, \alpha, \beta] = \frac{\alpha}{\alpha+\beta+n}$ and a posterior variance of $\mathbb{V}[\theta|Y = 0, \alpha, \beta] = \frac{\alpha(\beta+n)}{(\alpha+\beta+n)^2(\alpha+\beta+n+1)}$ and $\mathbb{E}[\theta|Y = n, \alpha, \beta] = \frac{\alpha+n}{\alpha+\beta+n}$ and $\mathbb{V}[\theta|Y = n, \alpha, \beta] = \frac{(\alpha+n)\beta}{(\alpha+\beta+n)^2(\alpha+\beta+n+1)}$. Note, the posterior variances are equivalent at these thresholds for all $\alpha = \beta$.

Parameters were given an uninformative Beta(1, 1) prior, which is identical to a standard uniform (Uniform(0, 1)) prior since it places equal probability on all parameter values within its support. This distribution has an expected value of $\mu = \frac{1}{2} = 0.500$ and a variance of $\sigma^2 = \frac{1}{12} \approx 0.083$. Further, the posterior is Beta($Y+1, n-Y+1$), from which various moments such as the expected value $\mathbb{E}[\theta|Y] = \frac{Y+1}{n+2}$ and variance $\mathbb{V}[\theta|Y] = \frac{(Y+1)(n-Y+1)}{(n+2)^2(n+3)}$, and other quantities, can be easily calculated. Clearly, $\mathbb{E}[\theta|Y = 0] = \frac{1}{n+2}$ and $\mathbb{V}[\theta|Y = 0] = \frac{n+1}{(n+2)^2(n+3)}$, and $\mathbb{E}[\theta|Y = n] = \frac{n+1}{n+2}$ and $\mathbb{V}[\theta|Y = n] = \frac{n+1}{(n+2)^2(n+3)}$. For example, with $n = 2$ sampled specimens for a given species (the minimum sample size possible) having no overlapping

379 genetic distances, on average, θ will be equal to $\mu = \frac{1}{4} = 0.250$, with a variance of
 380 $\sigma^2 = \frac{3}{80} \approx 0.0375$. In contrast, with $n = 100$ individuals and no overlap, the mean reduces to
 381 $\mu = \frac{1}{102} \approx 0.001$, with a variance of $\sigma^2 = 9.425 \times 10^{-5}$. When $n = 2$ and both taxon distances
 382 overlap, the mean is $\mu = \frac{3}{4} = 0.750$, having a variance of $\sigma^2 \approx 0.0375$. For $n = 100$ where all
 383 distances overlap, θ has expectation $\mu = \frac{101}{102} \approx 0.990$ and variance $\sigma^2 \approx 9.425 \times 10^{-5}$. These
 384 edge case examples demonstrate that the Beta(1, 1) distribution serves as a good first-pass
 385 prior for DNA barcode gap estimation using Phillips et al.'s (2024) metrics, though tighter
 386 bounds can likely be achieved. In general however, when possible, it is always advisable
 387 to incorporate prior information, even if only weak, rather than simply imposing complete
 388 ignorance in the form of a flat prior distribution. In the case of unimodal distributions, the
 389 (estimated) posterior mean often possesses the property that it readily decomposes into a
 390 convex linear combination, in the form of a weighted sum, of the (estimated) prior mean
 391 and the MLE. That is $\hat{\mu}_{\text{posterior}} = w\hat{\mu}_{\text{prior}} + (1 - w)\hat{\mu}_{\text{MLE}}$, where for the beta distribution,
 392 $w = \frac{\alpha + \beta}{\alpha + \beta + n}$. Therefore, with sufficient data, $w \rightarrow 0$ as $n \rightarrow \infty$, regardless of the values of α
 393 and β , and the choice of prior distribution becomes less important since the posterior will be
 394 dominated by the likelihood. For the Beta(1, 1), $w = \frac{2}{2+n}$, with $n = 2$ giving $w = \frac{1}{2} = 0.500$;
 395 that is, the posterior is the arithmetic average of the prior and the likelihood. For $n = 100$,
 396 the prior weight is $w = \frac{2}{102} \approx 0.0196$. The full Bayesian model for species x is thus given by

$$\begin{aligned}
 y_{\text{lwr}} &\sim \text{Binomial}(N, p_{\text{lwr}}) \\
 y_{\text{upr}} &\sim \text{Binomial}(M, p_{\text{upr}}) \\
 y'_{\text{lwr}} &\sim \text{Binomial}(N, p'_{\text{lwr}}) \\
 y'_{\text{upr}} &\sim \text{Binomial}(C, p'_{\text{upr}}) \\
 p_{\text{lwr}}, p_{\text{upr}}, p'_{\text{lwr}}, p'_{\text{upr}} &\sim \text{Beta}(1, 1).
 \end{aligned} \tag{12}$$

397 Note that p_x , q_x , p'_x , and q'_x in Equations (1)-(4) are denoted p_{lwr} , p_{upr} , p'_{lwr} , and q'_{upr} ,

respectively within Equation (12) for easy distinction between MLEs and Bayesian posterior estimates. The above statistical theory and derivations lay a good foundation for the remainder of this paper.

The proposed model is inherently vectorized to allow processing of multiple species datasets simultaneously. Model fitting was achieved using the Stan probabilistic programming language (Carpenter et al., 2017) framework for Hamiltonian Monte Carlo (HMC) (Neal, 2011) via the No-U-Turn Sampler (NUTS) sampling algorithm (Hoffman and Gelman, 2014) through the `rstan` R package (version 2.32.6) (Stan Development Team, 2023) in R (version 4.4.1) (R Core Team, 2024). Four Markov chains were run for 2000 iterations each in parallel across four cores with random parameter initializations. Within each chain, a total of 1000 samples was discarded as warmup (*i.e.*, burnin) to reduce dependence on starting conditions and to ensure posterior samples are reflective of the equilibrium distribution. Further, 1000 post-warmup draws were utilized per chain during the sampling phase. Because HMC/NUTS results in dependent samples that are minimally autocorrelated, chain thinning is not required. Each of these tuning parameters reflect default Markov Chain Monte Carlo (MCMC) settings in Stan to control both bias and variance, respectively, in the resulting draws. All analyses in the present work were carried out on a 2023 Apple MacBook Pro with M2 chip and 16 GB RAM running macOS Ventura 13.2. A random seed was set to ensure reproducibility of model results. Outputted estimates were rounded to three decimal places of precision. Posterior distributions were visualized as KDE plots using the `ggplot()` function within the `ggplot2` R package (version 3.5.1) (Wickham, 2016) with the default Gaussian kernel and optimal smoothness selection. To successfully run the Stan program, end users must have installed an appropriate compiler (such as GCC or Clang) which is compatible with their operating system, such as macOS.

Convergence was assessed both visually and quantitatively as follows: (1) through examining parameter traceplots, which depict the trajectory of accepted MCMC draws as a function of the number of iterations, (2) through monitoring the Gelman-Rubin potential

scale reduction factor statistic ($\hat{R}$) (Gelman and Rubin, 1992; Vehtari et al., 2021), which measures the concordance of within-chain *versus* between-chain variance, and (3) through calculating the effective sample size (ESS) for each parameter, which quantifies the number of independent samples generated Markov chains are equivalent to. Mixing of chains was deemed sufficient when traceplots looked like “fuzzy caterpillars”, $\hat{R} < 1.01$, and effective sample sizes were reasonably large (Gelman et al., 2020). After sampling, a number of summary quantities were reported, including posterior means, posterior SDs, and posterior quantiles from which 95% CrIs could be computed to make probabilistic inferences concerning true population parameters. To validate the overall correctness of the proposed statistical model given by Equation (12), as a means of comparison, posterior predictive checks (PPCs) were also employed to generate binomial random variates in the form of counts from the posterior predictive distribution; that is $\gamma = \{Np_x, Mq_x, Np'_x, Cq'_x\}$ to verify that the model adequately captures relevant features of the observed data. The proposed Bayesian model outlined here has a straightforward interpretation (**Table 1**).

3 Results and Discussion

The *Agabus* CYTB dataset analyzed by Phillips et al. (2024) is revisited herein for the species *A. bipustulatus* and *A. nevadensis*, since these taxa were the sole representatives for this locus, with the most and the least specimen records, respectively ($N = 701$ and $N = 2$) across all three assessed molecular markers. Briefly, using the R package **MACER** (Young et al., 2021), DNA sequences were downloaded from GenBank and BOLD and processed using a stringent workflow to obtain a 343-bp high-quality FASTA alignment representing 46 unique haplotypes. The resulting alignment contained no stop codons, ambiguous nucleotides, or other problematic artifacts that bias genetic diversity estimates. Further, both species- and genus-level distance outliers were excluded. Genetic distances were calculated using uncorrected p-distances. Further, *A. bipustulatus* comprised 45 total

Table 1: Interpretation of the DNA barcode gap estimators within $[a/a', b]$

Parameter	Explanation
p_x/p_{lwr}	When p_x/p_{lwr} is close to 0 (1), it suggests that the probability of intraspecific (interspecific) distances being larger (smaller) than interspecific (intraspecific) distances is low (high) on average, while the probability of interspecific (intraspecific) distances being larger (smaller) than intraspecific (interspecific) distances is high (low) on average; that is, there is (no) evidence for a DNA barcode gap.
q_x/p_{upr}	When q_x/p_{upr} is close to 0 (1), it suggests that the probability of interspecific (intraspecific) distances being larger (smaller) than intraspecific (interspecific) distances is high (low) on average, while the probability of intraspecific (interspecific) distances being larger (smaller) than interspecific (intraspecific) distances is low (high) on average; that is, there is (no) evidence for a DNA barcode gap.
p'_x/p'_{lwr}	When p'_x/p'_{lwr} is close to 0 (1), it suggests that the probability of intraspecific (combined interspecific distances for a target species and its nearest neighbour species) distances being larger (smaller) than combined interspecific distances for a target species and its nearest neighbour species (intraspecific distances) is low (high) on average, while the probability of combined interspecific distances for a target species and its nearest neighbour species (intraspecific distances) being larger (smaller) than intraspecific distances (combined interspecific distances for a target species and its nearest neighbour species) is high (low) on average; that is, there is (no) evidence for a DNA barcode gap.
q'_x/p'_{upr}	When q'_x/p'_{upr} is close to 0 (1), it suggests that the probability of combined interspecific distances for a target species and its nearest neighbour species (intraspecific distances) being larger (smaller) than intraspecific distances (combined interspecific distances for a target species and its nearest neighbour species) is high (low) on average, while the probability of intraspecific distances (combined interspecific distances for a target species and its nearest neighbour species) being larger (smaller) than combined interspecific distances for a target species and its nearest neighbour species (intraspecific distances) is low (high) on average; that is, there is (no) evidence for a DNA barcode gap.

haplotypes, whereas *A. nevadensis* possessed only one haplotype. Note, DNA barcode gap estimation is only possible for species having at least two specimen records. This dataset is a prime illustrative example highlighting the issue of inadequate taxon sampling, which arises frequently in large-scale phylogenetic and phylogeographic studies, in several respects. First, from a statistical viewpoint, sample sizes reflect extremes in reliable parameter estimation. Second, from a DNA barcoding perspective, *Agabus* currently comprises about 200 extant Linnean species according to the Global Biodiversity Information Facility (GBIF) (<https://www.gbif.org>); yet, due to the level of convenience sampling inherent in taxonomic collection efforts for this genus, adequate representation of species and genetic diversity is far from complete. Indeed, herein, species coverage for CYTB is only around 1.000%.

ECDF/CECDF plots were generated using the base R function `ecdf()`, along with `ggplot()`, and are displayed in **Figure 2**. Interpretation is simple: in the case of *A. bipustulatus*, p_x , all distances are greater than $a = 0$, with 36.900% of distances equal to a . Conversely, both q_x and q'_x equal one because all interspecific distances are less than $b \approx 0.026$; for *A. nevadensis*, both metrics are approximately equal to 0.010 since only about 1% of distances fall below $b = 0$.




Figure 2: ECDF/CECDF plots of the DNA barcode gap metrics for *A. bipustulatus* and *A. nevadensis* with a/a' and b clearly indicated as dashed vertical lines. When all genetic distances are equal to zero, no step function is shown.

MCMC parameter traceplots showed rapid mixing of chains to the stationary distribution

467 (Figure 3). Further, all $\hat{R}$ and ESS values (not shown) were close to their recommended
 468 cutoffs of one and thousands of samples, respectively, indicating chains are both well-mixed
 469 and have converged to the posterior distribution. Bayesian posterior estimates were reported
 470 alongside frequentist MLEs, in addition to SEs, posterior SDs, 95% CIs, and 95% CrIs
 471 (Table 2).




Figure 3: MCMC parameter traceplots applied to *A. bipustulatus* ([1]; $N = 701$) and *A. nevadensis* ([2]; $N = 2$) for CYTB across 1000 post-warmup iterations.

Table 2: Nonparametric frequentist and Bayesian estimates of distance distribution overlap/separation for the DNA barcode gap coalescent model parameters applied to *A. bipustulatus* ($N = 701$) and *A. nevadensis* ($N = 2$) for CYTB, including 95% CIs and CrIs. CrIs are based on 4000 posterior draws. All parameter estimates are reported to three decimal places of precision.

Species	Parameter	MLE (SE, 95% CI)	Bayes Est. (SD; 95% CrI)
<i>A. bipustulatus</i>	p_x/p_{lwr}	1.000 (0.000; 1.000-1.000)	1.000 (0.000; 1.000-1.000)
<i>A. bipustulatus</i>	q_x/p_{upr}	1.000 (0.000; 1.000-1.000)	1.000 (0.000; 0.999-1.000)
<i>A. bipustulatus</i>	p'_x/p'_{lwr}	1.000 (0.000; 1.000-1.000)	1.000 (0.000; 1.000-1.000)
<i>A. bipustulatus</i>	q'_x/p'_{upr}	1.000 (0.000; 1.000-1.000)	1.000 (0.000; 0.999-1.000)
<i>A. nevadensis</i>	p_x/p_{lwr}	1.000 (0.000; 1.000-1.000)	0.835 (0.144; 0.470-0.996)
<i>A. nevadensis</i>	q_x/p_{upr}	0.010 (0.002; 0.006-0.014)	0.010 (0.002; 0.007-0.014)
<i>A. nevadensis</i>	p'_x/p'_{lwr}	1.000 (0.000; 1.000-1.000)	0.834 (0.138; 0.481-0.994)
<i>A. nevadensis</i>	q'_x/p'_{upr}	0.010 (0.070; -0.128-0.148)	0.010 (0.002; 0.007-0.014)

CIs were calculated using the usual large sample $(1 - \alpha)100\%$ -level interval estimate given by $\hat{p} \pm z_{1-\frac{\alpha}{2}} \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$, where $z_{1-\frac{\alpha}{2}} = 1.960$ is the critical value for 95% confidence (*i.e.*, the 97.5th percent quantile from the standard Normal distribution), and α is the stated significance level (here, 5%). Given a $(1 - \alpha)100\%$ CI, with repeated sampling, on average $(1 - \alpha)100\%$ of constructed intervals will contain the true parameter of interest, assuming a correctly-specified statistical model; on the other hand, any given CI will either capture or exclude the true parameter with 100% certainty. This in stark contrast to a CrI, where the true parameter is contained within said interval with $(1 - \alpha)100\%$ probability, contingent on the data generating model being correct. Note, by default Stan computes equal-tailed (central) CrIs such that there is equal area situated in the left and right tails of the posterior distribution. For a 95% CrI, this corresponds to the 2.5th and 97.5th percent quantiles. However, constructed intervals are usually only valid for symmetric or nearly symmetric distributions. Given the bounded nature of the DNA barcode gap metrics, whose posterior distributions, as expected, show considerable skewness, an alternative approach to reporting CrIs, such as Highest Posterior Density (HPD) intervals (Chen and Shao, 1999) or shortest probability intervals (SPIn) (Liu et al., 2015) is warranted. As such asymmetric intervals generally attain greater statistical efficiency (in the form of smaller Mean Squared Error (MSE) or variance) and higher coverage probabilities than more standard interval estimates,

careful in-depth comparison is left for future work.

Findings based on nonparametric MLEs and Bayesian posterior means were quite comparable with one another and show evidence of complete overlap in intraspecific, interspecific, and combined interspecific distances for *A. bipustulatus* in both the $p/q/p_{\text{LWR}}/p_{\text{UPR}}$ and $p'/q'/p'_{\text{LWR}}/p'_{\text{UPR}}$ directions since the metrics attain magnitudes very close to one, with minimal noise (**Table 2**). As a result, this likely indicates that no DNA barcode gap is present for this species. Such findings are strongly reinforced by the very tight clustering of posterior draws (**Figure 4**) and associated interval estimates owing to the large number of specimens sampled for this species.




Figure 4: Scatterplots (black solid points) and distributions (black solid lines) depicting the DNA barcode gap metrics for *A. bipustulatus* ($N = 701$) and *A. nevadensis* ($N = 2$) across CYTB based on 4000 Bayesian posterior draws. MLEs and posterior means are displayed as coloured (red/blue) solid and dashed lines for the metrics, respectively.

On the other hand, the situation for *A. nevadensis* is more nuanced, as posterior values are further spread out (**Table 2** and **Figure 4**), suggesting less overall certainty in true parameter values given the low specimen sampling coverage for this taxon. Of note, calculated SEs and

SDs are large, with the 95% CIs and 95% CrIs being quite wide in most cases, consistent with much uncertainty regarding the computed frequentist and Bayesian posterior means of the DNA barcode gap metrics. For instance, the Bayesian analysis suggests that the data are consistent with both p_{lwr} and p'_{lwr} ranging from approximately 0.500-1.000 due to the large SD of about 0.140. The posterior means associated with these estimates themselves are also far from one. The CrI for p_{upr} and p'_{upr} also spans an order of magnitude, with the lower tail estimate of 0.007 being very close to zero. Further, regarding the frequentist analysis for the same species, the 95% CI for q_x is quite wide, reflecting considerable uncertainty in its true parameter value. Similarly, that for q'_x extends to negative values at the left endpoint, due to the corresponding SE of 0.070 being too high as a result of the extremely low sample size of $n = 2$ individuals sampled (**Table 2**). Since the 95% CI for q'_x truncated at the lower endpoint includes the value of zero, the null hypothesis for the presence of a DNA barcode gap cannot be rejected. Despite this, it is worth noting that truncation is not standard statistical practice and will likely lead to an interval with less than 95% nominal coverage having high relative error. In such cases, more appropriate confidence interval methods like the Wilson score interval, the exact (Clopper-Pearson) interval, or the Agresti-Coull interval should be employed (Newcombe, 1998; Agresti and Coull, 1998). KDEs for *A. bipustulatus* are strongly left (negatively) skewed (**Figure 4**), whereas those for *A. nevadensis* exhibit more symmetry, especially for p_{upr} and p'_{upr} (**Figure 4**). These differences are likely due to the stark contrast in sample sizes for the two examined species. Nevertheless, simulated counts of overlapping specimen records from the posterior predictive distribution (**Table 3**) were found to be very close to observed counts for both species, indicating that the proposed model adequately captures underlying variation.

Table 3: Posterior predictive checks of distance distribution overlap/separation for the DNA barcode gap coalescent model parameters applied to *A. bipustulatus* ($N = 701$) and *A. nevadensis* ($N = 2$) for CYTB. CrIs are based on 4000 posterior draws. All parameter estimates are reported to three decimal places of precision.

Species	Variable	Y	n	Bayes Est. (SD; 95% CrI)
<i>A. bipustulatus</i>	y_{lwr}	491401.000	491401.000	491400.018 (1.378; 491396.000-491401.000)
<i>A. bipustulatus</i>	y_{upr}	2804.000	2804.000	2803.019 (1.433; 2799.000-2804.000)
<i>A. bipustulatus</i>	y'_{jwr}	491401.000	491401.000	491400.008 (1.412; 491396.000-491401.000)
<i>A. bipustulatus</i>	y_{upr}	2804.000	2804.000	2802.992 (1.429; 2799.000-2804.000)
<i>A. nevadensis</i>	y_{lwr}	4.000	4.000	3.355 (0.888; 1.000-4.000)
<i>A. nevadensis</i>	y_{upr}	28.000	2804.000	29.151 (7.620; 16.000-45.000)
<i>A. nevadensis</i>	y'_{jwr}	4.000	4.000	3.325 (0.884; 1.000-4.000)
<i>A. nevadensis</i>	y_{upr}	28.000	2804.000	28.942 (7.409; 15.000-45.000)

Obtained results suggest that use of the Beta(1, 1) prior may not be appropriate given a low number of collected individuals for most taxa in DNA barcoding efforts. This suggests that further consideration of more informative beta priors is worthwhile. This is clear for *A. nevadensis*, where use of a stronger prior will likely eliminate the issue regarding accurate estimation of p_{lwr} and p'_{lwr} and stabilize their uncertainty.

4 Conclusion

Herein, the accuracy of the DNA barcode gap was analyzed from a rigorous statistical lens to expedite both the curation and growth of reference sequence libraries, ensuring they are populated with high quality, statistically defensible specimen records fit for purpose to address standing questions in ecology, evolutionary biology, management, and conservation. To accomplish this, recently proposed, easy to calculate, nonparametric MLEs were formally derived using ECDFs/CECDFs and applied to assess the extent of overlap/separation of distance distributions within and among two species of predatory water beetles in the genus *Agabus* sequenced at CYTB using a Bayesian binomial count model with conjugate beta priors. Findings highlight a high level of parameter uncertainty for *A. nevadensis*, whereas posterior estimates of the DNA barcode gap metrics for *A. bipustulatus* are much more certain. Based on these results, it is imperative that specimen sampling be prioritized to better reflect actual species boundaries. More generally, apart from the metrics being

employed to better highlighting the importance of within-species genetic diversity versus between-species divergence, it is expected that the approach developed herein will be of broad utility in applied fields, such as DNA-based detection of seafood fraud within global supply chains, and in the determination of species occupancy/detection probabilities at ecological sites of interest using active and passive environmental DNA (eDNA) methods such as metabarcoding.

Since the DNA barcode gap metrics often attain values very close to zero (suggesting no overlap and complete separation of distance distributions) and/or very near one (indicating no separation and complete overlap), in addition to more intermediate values, a noninformative $\text{Beta}(\frac{1}{2}, \frac{1}{2})$ prior may be more appropriate over complete ignorance imposed by a $\text{Beta}(1, 1)$ prior. The former distribution is U-shaped symmetric and places greater probability density at the extremes of the distribution due to its heavier tails, while still allowing for variability in parameter estimates within intermediate values along its domain. Note that this prior is Jeffreys' prior density (Jeffreys, 1946), which is proportional to the square root of the Fisher information $\mathcal{I}(\theta)$; that is, $\pi(\theta) \propto \theta^{-\frac{1}{2}}(1 - \theta)^{-\frac{1}{2}}$. Jeffreys' prior has several desirable statistical properties as a prior: that it is inversely proportional to the standard deviation of the binomial distribution, and most notably, that it is invariant to model reparameterization (Gelman et al., 2014). However, this prior can lead to divergent transitions, among other pathologies, imposed by complex geometry (*i.e.*, curvature) in the posterior space since many iterative stochastic MCMC sampling algorithms experience difficulties when exploring high density distribution regions. Thus, remedies to resolve them, such as lowering the step size of the HMC/NUTS sampler, should be attempted in future work, along with other approaches such as empirical Bayes estimation to approximate beta prior hyperparameters from observed data through the MLE or other methods of parameter estimation, such as the method of moments. Alternatively, hierarchical modelling could be employed to estimate separate distribution model hyperparameters for each species and/or compute distinct estimates for the directionality/comparison level of the DNA barcode gap metrics (*i.e.*, lower *vs.* upper,

570 non-prime *vs.* prime) separately within the genus under study. This would permit greater
571 flexibility through incorporating more fine-grained structure seen in the data; however, low
572 taxon sample sizes may preclude valid inferences to be reasonably ascertained due to the
573 large number additional parameters which would be introduced through the specification
574 of the hyperprior distributions. Methods outlined in Gelman et al. (2014, 2020), such as
575 dealing with non-exchangeability of observations and alternate model parameterizations like
576 the logit, may prove useful in this regard. Perhaps, the best course of action though is
577 the elicitation of priors from domain experts. Even though more work remains, it is clear
578 that both frequentist and Bayesian inference hold much promise for the future of molecular
579 biodiversity science.

Supplementary Information

None declared

Data Availability Statement

Raw data, R, and Stan code can be accessed via Dryad at:

<http://datadryad.org/stash/share/>

RZIfMixcEODe0RWP7eyXWQewSVbqEIA9UTrH3ZVKyn4.

A GitHub repository can be found at:

<https://github.com/jphill01/Bayesian-DNA-Barcode-Gap-Coalescent>.

Acknowledgements

We wish to recognise the valuable comments and discussions of Daniel (Dan) Gillis, Robert (Bob) Hanner, Robert (Rob) Young, and XXX anonymous reviewers.

We acknowledge that the University of Guelph resides on the ancestral lands of the Attawandaron people and the treaty lands and territory of the Mississaugas of the Credit. We recognize the significance of the Dish with One Spoon Covenant to this land and offer our respect to our Anishinaabe, Haudenosaunee and Métis neighbours as we strive to strengthen our relationships with them.

Funding

None declared.

Conflict of Interest

None declared.

Author Contributions

JDP wrote the manuscript, wrote R and Stan code, as well as analyzed and interpreted all model results.

References

Agresti, A. and B. A. Coull

1998. Approximate is better than ‘exact’ for interval estimation of binomial proportions. *The American Statistician*, 52(2):119–126.

Ahrens, D., F. Fujisawa, H.-J. Krammer, J. Eberle, S. Fabrizi, and A. Vogler

2016. Rarity and incomplete sampling in DNA-based species delimitation. *Systematic Biology*, 65(3):478–494.

Avise, J., J. Arnold, R. Ball, Jr., E. Bermingham, T. Lamb, J. Neigel, C. Reeb, and N. Saunders

1987. Intraspecific phylogeography: The mitochondrial DNA bridge between population genetics and systematics. *Annu. Rev. Ecol. Syst.*, 18:489–522.

Bartlett, S. and W. Davidson

1992. FINS (forensically informative nucleotide sequencing): A procedure for identifying the animal origin of biological specimens. *BioTechniques*, 12(3):408–411.

Bergsten, J., D. Bilton, T. Fujisawa, M. Elliott, M. Monaghan, M. Balke, L. Hendrich,

J. Geijer, J. Herrmann, G. Foster, I. Ribera, A. Nilsson, T. Barraclough, and A. Vogler

2012. The effect of geographical scale of sampling on DNA barcoding. *Systematic Biology*, 61(5):851–869.

Bouckaert, R., T. G. Vaughan, J. Barido-Sottani, S. Duchêne, M. Fourment,

A. Gavryushkina, J. Heled, G. Jones, D. Kühnert, N. De Maio, M. Matschiner, F. K.

Mendes, N. F. Müller, H. A. Ogilvie, L. du Plessis, A. Poppinga, A. Rambaut, D. Rasmussen,
I. Siveroni, M. A. Suchard, C.-H. Wu, D. Xie, C. Zhang, T. Stadler, and A. J. Drummond
2019. BEAST 2.5: An advanced software platform for Bayesian evolutionary analysis.
PLOS Computational Biology, 15(4):1–28.

Bucklin, A., D. Steinke, and L. Blanco-Bercial
2011. Dna barcoding of marine metazoa. *Annual Review of Marine Science*, 3(Volume 3,
2011):471–508.

Čandek, K. and M. Kuntner
2015. DNA barcoding gap: Reliable species identification over morphological and
geographical scales. *Molecular Ecology Resources*, 15(2):268–277.

Carpenter, B., A. Gelman, M. Hoffman, D. Lee, B. Goodrich, M. Betancourt, M. Brubaker,
J. Guo, P. Li, and A. Riddell
2017. Stan: A probabilistic programming language. *Journal of Statistical Software*, 76:1.

Carstens, B. C., T. A. Pelletier, N. M. Reid, and J. D. Satler
2013. How to fail at species delimitation. *Molecular Ecology*, 22(17):4369–4383.

Chen, M.-H. and Q.-M. Shao
1999. Monte Carlo estimation of Bayesian credible and HPD intervals. *Journal of
Computational and Graphical Statistics*, 8(1):69–92.

Collins, R. A. and R. H. Cruickshank
2013. The seven deadly sins of DNA barcoding. *Molecular Ecology Resources*,
13(6):969–975.

Collins, R. A. and R. H. Cruickshank
2014. Known knowns, known unknowns, unknown unknowns and unknown knowns in dna
barcoding: A comment on dowton et al. *Systematic Biology*, 63(6):1005–1009.

647 Dayrat, B.
 648 2005. Towards integrative taxonomy. *Biological Journal of the Linnean Society*,
 649 85(3):407–417.

650 de Queiroz, K.
 651 2007. Species concepts and species delimitation. *Systematic Biology*, 56(6):879–886.

652 Dempster, A. P., N. M. Laird, and D. B. Rubin
 653 1977. Maximum likelihood from incomplete data via the em algorithm. *Journal of the*
 654 *Royal Statistical Society: Series B (Methodological)*, 39(1):1–22.

655 Dowton, M., K. Meiklejohn, S. L. Cameron, and J. Wallman
 656 2014. A preliminary framework for DNA barcoding, incorporating the multispecies
 657 coalescent. *Systematic Biology*, 63(4):639–644.

658 Efron, B. and R. J. Tibshirani
 659 1993. *An Introduction to the Bootstrap*. Chapman & Hall.

660 Flouri, T., X. Jiao, B. Rannala, and Z. Yang
 661 2018. Species tree inference with BPP using genomic sequences and the multispecies
 662 coalescent. *Molecular Biology and Evolution*, 35(10):2585–2593.

663 Fonseca, E. M., D. J. Duckett, and B. C. Carstens
 664 2021. P2C2M.GMYC: An R package for assessing the utility of the Generalized Mixed
 665 Yule Coalescent model. *Methods in Ecology and Evolution*, 12(3):487–493.

666 Fujisawa, T. and T. Barraclough
 667 2013. Delimiting species using single-locus data and the generalized mixed yule coalescent
 668 approach: A revised method and evaluation on simulated data sets. *Systematic Biology*,
 669 62(5):707–724.

670 Fujita, M. K., A. D. Leaché, F. T. Burbrink, J. A. McGuire, and C. Moritz
671 2012. Coalescent-based species delimitation in an integrative taxonomy. *Trends in Ecology*
672 *& Evolution*, 27(9):480–488.

673 Gelman, A., J. Carlin, H. Stern, D. Duncan, A. Vehtari, and D. Rubin
674 2014. *Bayesian Data Analysis*, third edition. Chapman and Hall/CRC.

675 Gelman, A. and D. Rubin
676 1992. Inference from iterative simulation using multiple sequences. *Statistical Science*,
677 7(4):457–472.

678 Gelman, A., A. Vehtari, D. Simpson, C. Margossian, B. Carpenter, Y. Yao, L. Kennedy,
679 J. Gabry, P.-C. Bürkner, and M. Modrák
680 2020. Bayesian workflow.

681 Hart, M. and J. Sunday
682 2007. Things fall apart: biological species form unconnected parsimony networks. *Biology*
683 *Letters*, 3(5):509–512.

684 Hebert, P., A. Cywinska, S. Ball, and J. deWaard
685 2003a. Biological identifications through DNA barcodes. *Proceedings of the Royal Society*
686 *of London B: Biological Sciences*, 270(1512):313–321.

687 Hebert, P., S. Ratnasingham, and J. de Waard
688 2003b. Barcoding animal life: Cytochrome c oxidase subunit 1 divergences among
689 closely related species. *Proceedings of the Royal Society of London B: Biological Sciences*,
690 270(Suppl 1):S96–S99.

691 Hebert, P. D., M. Y. Stoeckle, T. S. Zemlak, and C. M. Francis
692 2004. Identification of birds through DNA barcodes. *PLoS Biol*, 2(10):e312.

693 Hickerson, M. J., C. P. Meyer, and C. Moritz
694 2006. DNA barcoding will often fail to discover new animal species over broad parameter
695 space. *Systematic Biology*, 55(5):729–739.

696 Hoffman, M. and A. Gelman
697 2014. The No-U-Turn Sampler: Adaptively setting path lengths in Hamiltonian Monte
698 Carlo. *Journal of Machine Learning Research*, 15:1593–1623.

699 Hubert, N. and R. Hanner
700 2015. DNA barcoding, species delineation and taxonomy: A historical perspective. *DNA*
701 *Barcodes*, 3:44–58.

702 Hubert, N., J. Phillips, and R. Hanner
703 2024. Delimiting species with single-locus dna sequences. In *DNA Barcoding*, R. DeSalle,
704 ed., volume 2744 of *Methods in Molecular Biology*. New York, NY: Humana.

705 Jackson, N. D., B. C. Carstens, A. E. Morales, and B. C. O’Meara
706 2017a. Species delimitation with gene flow. *Systematic Biology*, 66(5):799–812.

707 Jackson, N. D., A. E. Morales, B. C. Carstens, and B. C. O’Meara
708 2017b. PHRAPL: Phylogeographic inference using approximate likelihoods. *Systematic*
709 *Biology*, 66(6):1045–1053.

710 Jeffreys, H.
711 1946. An invariant form for the prior probability in estimation problems. *Proceedings*
712 *of the Royal Society of London. Series A, Mathematical and Physical Sciences*,
713 186(1007):453–461.

714 Jukes, T. and C. Cantor
715 1969. Evolution of protein molecules. In *Mammalian Protein Metabolism*, H. N. Munro,
716 ed., Pp. 21–132. New York: Academic Press.

717 Kapli, P., S. Lutteropp, J. Zhang, K. Kobert, P. Pavlidis, A. Stamatakis, and T. Flouri
718 2017. Multi-rate poisson tree processes for single-locus species delimitation under maximum
719 likelihood and markov chain monte carlo. *Bioinformatics*, 33(11):1630–1638.

720 Kimura, M.
721 1980. A simple method for estimating evolutionary rates of base substitutions
722 through comparative studies of nucleotide sequences. *Journal of Molecular Evolution*,
723 16(1):111–120.

724 Kingman, J.
725 1982a. The coalescent. *Stochastic Processes and Their Applications*, 13:235–248.

726 Kingman, J.
727 1982b. On the genealogy of large populations. *Journal of Applied Probability*, 19(A):27–43.

728 Knowles, L. L.
729 2009. Statistical phylogeography. *Annual Review of Ecology, Evolution, and Systematics*,
730 40(Volume 40, 2009):593–612.

731 Knowles, L. L. and W. P. Maddison
732 2002. Statistical phylogeography. *Molecular Ecology*, 11(12):2623–2635.

733 Kullback, S. and R. Leibler
734 1951. On information and sufficiency. *Annals of Mathematical Statistics*, 22(1):79–86.

735 Leaché, A., M. Fujita, V. Minin, and R. Bouckaert
736 2014. Species delimitation using genome-wide snp data. *Systematic Biology*, 63(4):534–542.

737 Liu, Y., A. Gelman, and T. Zheng
738 2015. Simulation-efficient shortest probability intervals. *Statistical Computing*, 25:809–819.

Mather, N., S. Traves, and S. Ho

2019. A practical introduction to sequentially Markovian coalescent methods for estimating demographic history from genomic data. *Ecology and Evolution*, 10(1):579–589.

Mayr, E.

1942. *Systematics and the Origin of Species from the Viewpoint of a Zoologist*. New York: Columbia University Press.

Meier, R., G. Zhang, and F. Ali

2008. The use of mean instead of smallest interspecific distances exaggerates the size of the “barcoding gap” and leads to misidentification. *Systematic Biology*, 57(5):809–813.

Meyer, C. and G. Paulay

2005. DNA barcoding: Error rates based on comprehensive sampling. *PLOS Biology*, 3(12):e422.

Monaghan, M., R. Wild, M. Elliot, T. Fujisawa, M. Balke, D. Inward, D. Lees, R. Ranaivosolo, P. Eggleton, T. Barraclough, and A. Vogler

2009. Accelerated species inventory on madagascar using coalescent-based models of species delineation. *Systematic Biology*, 58:298–311.

Mutanen, M., S. M. Kivelä, R. A. Vos, C. Doorenweerd, S. Ratnasingham, A. Hausmann, P. Huemer, V. Dincă, E. J. van Nieukerken, C. Lopez-Vaamonde, R. Vila, L. Aarvik, T. Decaëns, K. A. Efetov, P. D. N. Hebert, A. Johnsen, O. Karsholt, M. Pentinsaari, R. Rougerie, A. Segerer, G. Tarmann, R. Zahiri, and H. C. J. Godfray

2016. Species-level para- and polyphyly in DNA barcode gene trees: Strong operational bias in european lepidoptera. *Systematic Biology*, 65(6):1024–1040.

Neal, R.

2011. MCMC using Hamiltonian dynamics. In *Handbook of Markov Chain Monte Carlo*. Chapman & Hall, CRC Press.

Newcombe, R. G.

1998. Two-sided confidence intervals for the single proportion: comparison of seven methods. *Statistics in Medicine*, 17(8):857–872.

Pante, E., N. Puillandre, A. Viricel, S. Arnaud-Haond, D. Aurelle, M. Castelin, A. Chenuil, C. Destombe, D. Forcioli, M. Valero, F. Viard, and S. Samadi

2015. Species are hypotheses: Avoid connectivity assessments based on pillars of sand. *Molecular Ecology*, 24(3):525–544.

Paradis, E., J. Claude, and K. Strimmer

2004. Ape: Analyses of phylogenetics and evolution in r language. *Bioinformatics*, 20(2):289–290.

Pentinsaari, M., H. Salmela, M. Mutanen, and T. Roslin

2016. Molecular evolution of a widely-adopted taxonomic marker (COI) across the animal tree of life. *Scientific Reports*, 6:35275.

Phillips, J., T. Athey, P. McNicholas, and R. Hanner

2023. VLF: An r package for the analysis of very low frequency variants in DNA sequences. *Biodiversity Data Journal*, 11:e96480.

Phillips, J., D. Gillis, and R. Hanner

2019. Incomplete estimates of genetic diversity within species: Implications for DNA barcoding. *Ecology and Evolution*, 9(5):2996–3010.

Phillips, J., D. Gillis, and R. Hanner

2020. HACSIm: An R package to estimate intraspecific sample sizes for genetic diversity assessment using haplotype accumulation curves. *PeerJ Computer Science*.

Phillips, J., D. Gillis, and R. Hanner

2022. Lack of statistical rigor in DNA barcoding likely invalidates the presence of a true species’ barcode gap. *Frontiers in Ecology and Evolution*, 10:859099.

789 Phillips, J., C. Griswold, R. Young, N. Hubert, and H. Hanner
790 2024. *A Measure of the DNA Barcode Gap for Applied and Basic Research*, Pp. 375–390.
791 New York, NY: Springer US.

792 Phillips, J., R. Gwiazdowski, D. Ashlock, and R. Hanner
793 2015. An exploration of sufficient sampling effort to describe intraspecific DNA barcode
794 haplotype diversity: examples from the ray-finned fishes (Chordata: Actinopterygii). *DNA*
795 *Barcodes*, 3:66–73.

796 Pons, J., T. G. Barraclough, J. Gomez-Zurita, A. Cardoso, D. P. Duran, S. Hazell,
797 S. Kamoun, W. D. Sumlin, and A. P. Vogler
798 2006. Sequence-based species delimitation for the DNA taxonomy of undescribed insects.
799 *Systematic Biology*, 55(4):595–609.

800 Puillandre, N., S. Brouillet, and G. Achaz
801 2021. ASAP: Assemble Species by Automatic Partitioning. *Molecular Ecology Resources*,
802 21(2):609–620.

803 Puillandre, N., A. Lambert, S. Brouillet, and G. Achaz
804 2011. ABGD, Automatic Barcode Gap Discovery for primary species delimitation.
805 *Molecular Ecology*, 21(8):1864–1877.

806 R Core Team
807 2024. *R: A Language and Environment for Statistical Computing*. R Foundation for
808 Statistical Computing, Vienna, Austria.

809 Rannala, B.
810 2015. The art and science of species delimitation. *Current Zoology*, 61(5):846–853.

811 Rannala, B. and Z. Yang
812 2003. Bayes estimation of species divergence times and ancestral population sizes using
813 DNA sequences from multiple loci. *Genetics*, 164:1645–1656.

814 Rannala, B. and Z. Yang
815 2017. Efficient Bayesian species tree inference under the multispecies coalescent. *Systematic*
816 *Biology*, 66(5):823–842.

817 Ratnasingham, S. and P. Hebert
818 2007. BOLD: The Barcode of Life Data System (<http://www.barcodinglife.org>). *Molecular*
819 *Ecology Notes*, 7(3):355–364.

820 Ratnasingham, S. and P. D. N. Hebert
821 2013. A dna-based registry for all animal species: The barcode index number (bin) system.
822 *PLoS One*, 8(7):e66213.

823 Stan Development Team
824 2023. RStan: The R interface to Stan. R package version 2.32.6.

825 Stoeckle, M. and D. Thaler
826 2014. DNA barcoding works in practice but not in (neutral) theory. *PLoS One*,
827 9(7):e100755.

828 Talavera, G., V. Dincă, and R. Vila
829 2013. Factors affecting species delimitations with the GMYC model: insights from a
830 butterfly survey. *Methods in Ecology and Evolution*, 4(12):1101–1110.

831 Templeton, A., K. Crandall, and C. Sing
832 1992. A cladistic analysis of phenotypic associations with haplotypes inferred from
833 restriction endonuclease mapping and DNA sequence data. III. Cladogram estimation.
834 *Genetics*, 132(2):619–633.

835 Vehtari, A., A. Gelman, D. Simpson, B. Carpenter, and P.-C. Bürkner
836 2021. Rank-normalization, folding, and localization: An improved $\hat{R}$ for assessing
837 convergence of MCMC (with discussion). *Bayesian Analysis*, 16(2):667–718.

- Wells, T., T. Carruthers, P. Muñoz-Rodríguez, A. Sumadijaya, J. R. I. Wood, and R. W. Scotland
2022. Species as a heuristic: Reconciling theory and practice. *Systematic Biology*, 71(5):1233–1243.
- Wickham, H.
2016. *ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York.
- Wiemers, M. and K. Fiedler
2007. Does the DNA barcoding gap exist? - a case study in blue butterflies (Lepidoptera: Lycaenidae). *Frontiers in Zoology*, 4(8).
- Yang, Z. and B. Rannala
2010. Bayesian species delimitation using multilocus sequence data. *Proceedings of the National Academy of Sciences*, 107:9264–9269.
- Yang, Z. and B. Rannala
2014. Unguided species delimitation using dna sequence data from multiple loci. *Molecular Biology and Evolution*, 31(12):3125–3135.
- Yang, Z. and B. Rannala
2017. Bayesian species identification under the multispecies coalescent provides significant improvements to DNA barcoding analyses. *Molecular Ecology*, 26:3028–3036.
- Young, R., R. Gill, D. Gillis, and R. Hanner
2021. Molecular Acquisition, Cleaning and Evaluation in R (MACER) - A tool to assemble molecular marker datasets from BOLD and GenBank. *Biodiversity Data Journal*, 9:e71378.
- Zhang, J., P. Kapli, P. Pavlidis, and A. Stamatakis
2013. A general species delimitation method with applications to phylogenetic placements. *Bioinformatics*, 29(22):2869–2876.