

A Bayesian Model of the DNA Barcode Gap

Jarrett D. Phillips^{1,2*} (ORCID: 0000-0001-8390-386X),

Nicolas Hubert³ (ORCID: 0000-0001-9248-3377)

Robert H. Hanner² (ORCID: 0000-0003-0703-1646)

¹*School of Computer Science, University of Guelph, Guelph, ON., Canada, N1G2W1*

²*Department of Integrative Biology, University of Guelph, Guelph, ON., Canada, N1G2W1*

Université de Montpellier, Place Eugène Bataillon, 34095 Montpellier cedex 05, France

***Corresponding Author:** Jarrett D. Phillips¹

Email Address: jphill01@uoguelph.ca

Running Title: Bayesian inference for DNA barcode gap estimation

Abstract

A simple statistical model of the DNA barcode gap is outlined. Specifically, accuracy of recently introduced nonparametric metrics, inspired by coalescent theory, to characterize the extent of proportional overlap/separation in maximum and minimum pairwise genetic distances within and among species, respectively, is explored in both frequentist and Bayesian contexts. The empirical cumulative distribution function (ECDF) is utilized to estimate probabilities associated with positively skewed extreme tail distribution quantiles bounded on the closed unit interval $[0, 1]$ based on a straightforward binomial distance overlap count. Statistical moments, notably, the mean and variance, are derived for edge cases of no overlap/complete separation and full overlap/no separation in genomic sequence variation, and are shown to be quite tight as sample sizes increase. Further, a new way to visualize distances and the DNA barcode gap estimators appears revealing. Using R and the probabilistic programming language, Stan, the proposed maximum likelihood estimators and Bayesian model are demonstrated on cytochrome *b* (CYTB) gene sequences from two *Agabus* diving beetle species, *A. bispusulatus* and *A. nevadensis*, exhibiting limits in the extent of representative taxonomic sampling ($N = 701$ and $N = 2$ individuals, respectively). Large-sample theory and MCMC simulations clearly expose problems, showing much uncertainty in parameter estimates, particularly under the latter paradigm, and when specimen sample sizes for target species are small. Findings herein highlight the promise of the Bayesian approach using a conjugate beta prior for reliable posterior estimation over classical inference when available data are sparse. Obtained results can help shed light on foundational and applied research questions concerning DNA-based specimen identification and species delineation for studies in evolutionary biology and ecology, as well as biodiversity conservation and management, of wide-ranging taxa.

Keywords: Bayesian/frequentist inference, DNA barcoding, intraspecific genetic distance, interspecific genetic distance, specimen identification, species discovery

1 Introduction

The routine use of DNA sequences to support broad evolutionary hypotheses and questions concerning central demographic processes, like gene flow and speciation, is now ubiquitous. Since the late 1980s, mitochondrial DNA (mtDNA) has been employed to produce a distinctive and measurable pattern of genetic polymorphism in diverse and spatially-distributed taxonomic lineages such as birds, fishes, insects, and arachnids, among other extensively studied groups (Avise et al., 1987). The application of genomic data to applied fields like biodiversity forensics, conservation, and management for the molecular identification of unknown specimen samples came later (*e.g.*, Forensically Informative Nucleotide Sequencing (FINS); Bartlett and Davidson (1992)). Since its inception over 20 years ago, DNA barcoding (Hebert et al., 2003a,b) built significantly on earlier work and has emerged as a robust method of specimen identification and species discovery across myriad multicellular eukaryotes which have been sequenced at easily obtained short, standardized gene regions like the cytochrome *c* oxidase subunit I (5'-COI) mitochondrial locus for animals. However, DNA barcoding is not without its controversies.

The success of the single-locus approach, particularly for regulatory and forensic applications, depends crucially on two important factors: (1) the availability of high-quality specimen records found in public reference sequence databases such as the Barcode of Life Data Systems (BOLD; <http://www.barcodinglife.org>) (Ratnasingham and Hebert, 2007) and GenBank (<https://www.ncbi.nlm.nih.gov/genbank/>), and (2) the establishment of a DNA barcode gap — the notion that the maximum genetic distance observed within species is much smaller than the minimum degree of marker variation found among species (Meyer and Paulay, 2005; Meier et al., 2008). Early research has demonstrated that the presence of a DNA barcode gap hinges strongly on extant levels of species haplotype diversity gauged from comprehensive specimen sampling at wide geographic and ecological scales (Bergsten et al., 2012; Čandek and Kuntner, 2015). Despite this, many taxonomic groups lack adequate separation in their pairwise intraspecific and interspecific genetic distances due to varying

rates of evolution in both genes and taxa (Pentinsaari et al., 2016). Furthermore, it has been well-demonstrated that the presence of a DNA barcode gap becomes less certain with increasing spatial scale of sampling, since interspecific distances increase due to population structuring under limited dispersal through isolation-by-distance, while intraspecific distances shrink as more closely-related species are sampled (Phillips et al., 2022). This can pose problems in cases of rare singleton species or monotypic taxa, for instance (Ahrens et al., 2016). As a result, rapid matching of unknown samples to expertly-validated references can be compromised, leading to cases of false positives (taxon oversplitting) and false negatives (excessive lumping of taxa) due to incomplete lineage sorting, interspecies hybridization, genome introgression, species synonymy, cryptic species diversity, and misidentifications (Hubert and Hanner, 2015; Phillips et al., 2022). Another culprit includes PCR/sequencing error caused by incorrect nucleotide basecalling, chimeras/heteroplasmies, sample contamination, indels (insertion/deletions), NUMT (nuclear-mitochondrial insert) pseudogene amplification, or *Wolbachia* endosymbiont infection, among other explanations (Phillips et al., 2023). Each of these can artificially magnify levels of genetic diversity, rendering techniques like DNA barcoding unreliable for certain groups like marine invertebrates, which are both undersampled and possess slow rates of molecular evolution (Bucklin et al., 2011).

Recent work has argued that DNA barcoding, in its current form, is lacking in statistical rigor (Phillips et al., 2022). Most publications rely strongly on heuristic distance-based thresholds calculated from convenient specimen sample sizes (usually 5-10 specimens per species, but singleton/doubletons are more prevalent), rather than on optimal sampling designs using collated taxon haplotype information (Phillips et al., 2015, 2020, 2022), to infer taxonomic identity and delineate boundaries on the basis of clustering variation into putative groups reminiscent of species. Of these studies, few report measures of uncertainty, such as standard errors (SEs) and confidence intervals (CIs), around estimates of intraspecific and interspecific variation (*e.g.*, Wiemers and Fiedler (2007)), calling into question the existence

92 of a true species' DNA barcode gap (Čandek and Kuntner, 2015; Phillips et al., 2022).
 93 To support this idea, novel nonparametric locus- and species-specific metrics based on the
 94 multispecies coalescent (MSC) were recently outlined by Phillips et al. (2024). The basic
 95 coalescent (Kingman, 1982a,b) encompasses a backwards continuous-time stochastic Markov
 96 process of allelic sampling among lineages under the infinite sites model of mutation within
 97 a large neutrally-evolving, panmictic, diploid species population of constant size from the
 98 present into the past towards the most recent common ancestor (MRCA). Although the
 99 coalescent sampling process plays a foundational role in evolutionary biology and population
 100 genetics theory due to its inherent simplicity and flexibility (Rosenberg and Nordborg, 2002),
 101 the approach has been both under-utilized and under-appreciated as a key player within DNA
 102 barcoding (Collins and Cruickshank, 2014; Downton et al., 2014; Hubert and Hanner, 2015;
 103 Phillips et al., 2022). Numerous theoretical and empirical investigations have shown that
 104 protein-coding sequence variation within haploid loci like COI is generally subject to strong
 105 purifying selection and selective sweeps at nonsynonymous sites, rather than the occurrence
 106 of functionally silent substitutions within synonymous regions (Stoeckle and Thaler, 2014).
 107 However, because mitochondrial markers have considerably smaller effective population sizes
 108 (N_e) compared to nuclear regions of the genome, accumulated mutations become fixed,
 109 and therefore diagnostic of species, at a much faster rate due to genetic drift (Hubert and
 110 Hanner, 2015; Phillips et al., 2022). Previously proposed MSC algorithmic approaches (of
 111 which there are too many to exhaustively list here, and which have mostly been applied
 112 to multilocus sequence data), generally assume a strict molecular clock and a simplified
 113 model of DNA sequence evolution across closely-related taxa (Ho and Duchêne, 2014), from
 114 which an estimated species phylogeny may be constructed (*e.g.*, with or without use of a
 115 guide tree) (*e.g.*, Rannala and Yang (2003, 2017); Yang and Rannala (2010, 2014, 2017)).
 116 Even tree- and network-based taxon delimitation methods designed for single loci, such as
 117 the Generalized Mixed Yule Coalescent (GMYC) (Pons et al., 2006; Monaghan et al., 2009;
 118 Fujisawa and Barraclough, 2013), Poisson Tree Processes (PTP) (Zhang et al., 2013; Kapli

et al., 2017), and statistical parsimony/haplotype networks (*e.g.*, TCS networks (Templeton et al., 1992)), which have seen much use in DNA barcoding initiatives, have their flaws. Optimal performance of these methods relies heavily on careful parameter selection (*e.g.*, prior specification of branching rates in ultrametric trees in the case of GMYC) (Hubert et al., 2024) within third party software like BEAST (Bouckaert et al., 2019), among other concerns (Talavera et al., 2013; Fonseca et al., 2021). Hubert et al. (2024) categorized statistical species delimitation methods according to five crucial factors: availability, reliability, scalability, understandability, and usability. Any proposed approach should score high across all these aspects, yet most current methods fall short. (See Table 1 in Hubert et al. (2024) for a comprehensive overview.). Careless parameterization of established approaches by users unfamiliar with their underlying statistical characteristics may severely bias overall results in unintended ways (Hart and Sunday, 2007). In contrast, Phillips et al.’s (2024) approach is both tree- and network-free, and does not require judicious parameter setting; all that is needed is a distance matrix. Therefore, the method is easy to use and interpret, quite efficient, and very fast to run on large datasets. Most importantly, the method satisfies all criteria mentioned by Hubert et al. (2024).

The DNA barcode gap statistics have been shown to hold strong promise for reliable DNA barcode gap assessment when applied to predatory *Agabus* (Coleoptera: Dytiscidae) diving beetles (Phillips et al., 2024). With a Holarctic range, despite their ease of sampling and well-established taxonomy, this group possesses few morphologically-distinct taxonomic characters that readily facilitate their assignment to the species level (Bergsten et al., 2012). Accurate taxon discrimination within *Agabus* is paramount since several species are currently classified as endangered on the International Union for Conservation of Nature (IUCN) Red List (<https://www.iucnredlist.org>) due to habitat fragmentation/loss, water pollution, and climate change. Further, the proposed metrics indicate that sister species pairs from this taxon are often difficult to distinguish on the basis of their DNA barcode sequences (Phillips et al., 2024), which may be caused by *Agabus*’s increased degree of non-monophyly

at wider geospatial scales, occurrence of hybridization/introgression, or presence of synonymy (Bergsten et al., 2012). Using sequence data from three mitochondrial cytochrome markers (5'-COI, 3'-COI, and cytochrome *b* (CYTB)) obtained from BOLD and GenBank, results from Phillips et al. (2024) highlight that DNA barcoding has been a one-sided argument. Phillips et al.'s (2024) findings point to the need to balance both the sufficient collection of specimens, as well as the extensive sampling of species. Not surprisingly, DNA barcode libraries for most taxonomic groups are biased toward the latter effort (Phillips et al., 2015, 2019, 2022).

The estimators from Phillips et al. (2024) represent a clear improvement over simple, yet arbitrary, distance heuristics such as the 2% rule noted by Hebert et al. (2003a) and the 10 $\times$ rule (Hebert et al., 2004) that form the core of threshold-based single-locus species delimitation tools like Automatic Barcode Gap Discovery (ABGD) (Puillandre et al., 2011), Assemble Species by Automatic Partitioning (ASAP) (Puillandre et al., 2021), and the Barcode Index Number (BIN) framework (Ratnasingham and Hebert, 2013). The 2% rule asserts that DNA barcodes differing by at least 2% at sequenced genomic regions should be expected to originate from different biological species, whereas the 10 $\times$ rule suggests that barcode sequences displaying 10 times more genetic variation among species than within taxa is suggestive of a distinct evolutionary origin. However, the lack of adoption of an explicit and universally agreed-upon species concept in DNA barcoding studies that readily governs lineage formation and proliferation necessary to establish rigorous taxon definitions for successful delimitation of hypothesized and heuristic evolutionary units using these well-known criteria, in conjunction with secondary lines of evidence (*e.g.*, morphology, ecology, geography, ontogeny, and behaviour) promised by an integrative framework (Dayrat, 2005), is missing (de Queiroz, 2007; Rannala, 2015; Pante et al., 2015; Wells et al., 2022). This is the case, for instance, with species existing in allopatry/sympatry, where establishment of reproductive isolation on the basis of fixed diagnostic differences using the Biological Species Concept (BSC) (Mayr, 1942), is challenging (Fujita et al., 2012). In addition, the reliance on

visualization approaches, such as frequency histograms, dotplots, and quadrant plots, to
 expose DNA barcoding’s limitations, has also been criticized since they can fog real species
 signal relative to noise (Collins and Cruickshank, 2013; Phillips et al., 2022). Up until the
 work of Phillips et al. (2024), the majority of studies (*e.g.*, Young et al. (2021)) have treated
 the DNA barcode gap as a binary response. Given poor sampling depth for most taxa, a
 Yes/No dichotomy is inherently flawed because it can falsely imply a DNA barcode gap is
 present for a taxon of interest when in fact no such separation in genetic distances exists.
 On the other hand, in the case of sequence overlap, while the presence a single specimen
 record is a necessary requirement to nullify the existence of a DNA barcode gap in light
 of available data, it is not a sufficient condition, since said taxon record could represent a
 species misidentification or other error. Thus, the MSC represents an ideal model to settle
 ongoing debates and issues regarding DNA barcoding’s use as both a molecular identification
 and delineation tool. Notably, Fujita et al. (2012) remarks that the MSC applied to the
 species delimitation problem harmonizes many species concepts since it identifies independent
 evolutionary trajectories under the umbrella of the General Lineage Concept. This makes
 the MSC a fitting tool to model speciation and other modes of diversification (Collins and
 Cruickshank, 2014). The recently introduced DNA barcode gap metrics go some way into
 further operationalizing this task.

Phillips et al.’s (2024) proposed statistics quantify the extent of asymmetric directionality
 of proportional distance distribution overlap/separation for species within well-sampled
 taxonomic genera based on a straightforward distance count, in a similar vein to established
 measures of statistical similarity such as the Kullback-Leibler (KL) divergence (Kullback
 and Leibler, 1951) and other related statistics. The metrics can be employed in a variety
 of ways, including to validate performance of marker genes for specimen identification to
 the species level (using nonparametric bootstrapping, as in Phillips et al. (2024)). The
 estimators can further be utilized to assess whether computed values are consistent with
 fundamental population genetic-level parameters like N_e , mutation rates (μ), migration rates

(m) , and divergence times (τ), as well as derived ones, such as the population mutation rate
 $\theta = 4N_e\mu$, and the fixation index, F_{ST} , for species under study in a statistical phylogeographic
 setting (Knowles and Maddison, 2002; Knowles, 2009; Mather et al., 2019). Early on, species
 discovery using DNA barcoding was presumed to only work for reciprocally monophyletic
 groups, and thus concerned itself with terminal branches of generated phylogenies rather
 than more basal lineages occurring deeper in hypothesized species trees (Fujita et al., 2012;
 Mutanen et al., 2016). Furthermore, the occurrence of short branches and large N_e within
 resolved phylogenies increases the probability of deep coalescence, clouding species
 delimitations, which often fail or are uncertain in broad parameter space (Carstens et al.,
 2013; Hickerson et al., 2006; Rannala, 2015). As DNA barcoding is a single-locus approach,
 it is problematic for evolutionarily young taxa. Incomplete lineage sorting within gene
 genealogies, leading to the retention of ancestral polymorphisms within populations, is a
 common phenomenon due to the ongoing stochastic dynamic of mutation generating
 population variation, and genetic drift driving gene variants to fixation (Rannala, 2015). The
 most promising way forward in assessing the validity of Phillips et al.’s (2024) method seems
 to be through the use of software such as BPP (Bayesian Phylogenetics and Phylogeography),
 which permits efficient full Bayesian simulations under various MSC models (*e.g.*, MSC-I
 (MSC with introgression) or MSC-M (MSC with migration), among others) using MCMC for
 tree parameter estimation (via the A00 option, for instance) (Flouri et al., 2018), or PHRAPL
 (Phylogeographic Inference using Approximate Likelihoods) (Jackson et al., 2017a,b), which
 employs tractable phylogenetic likelihood calculations via the genealogical divergence index
 (gdi). Other hypothesis-driven approaches to species delimitation model selection, like Bayes
 Factors (BFs) (Leaché et al., 2014), are also possible.

While introduction of the new metrics is a step in the right direction, what appears to
 be missing is a rigorous statistical treatment of the DNA barcode gap. This includes an
 unbiased way to compute the statistical accuracy of Phillips et al.’s (2024) estimators arising
 through problems inherent in frequentist maximum likelihood estimation for probability

distributions having bounded positive support on the closed unit interval $[0, 1]$. To this end, here, a Bayesian model of the DNA barcode gap coalescent is introduced to rectify such issues. The model allows accurate estimation of posterior means, posterior standard deviations (SDs), posterior quantiles, and credible intervals (CrIs) for the statistics given datasets of intraspecific and interspecific distances for species of interest.

2 Methods

2.1 DNA Barcode Gap Metrics

The novel nonparametric maximum likelihood estimators (MLEs) of proportional overlap/separation between intraspecific and interspecific distance distributions for a given species (x) to aid assessment of the DNA barcode gap are as follows:

$$p_x = \frac{\#\{d_{ij} \geq a\}}{\#\{d_{ij}\}} \quad (1)$$

$$q_x = \frac{\#\{d_{XY} \leq b\}}{\#\{d_{XY}\}} \quad (2)$$

$$p'_x = \frac{\#\{d_{ij} \geq a'\}}{\#\{d_{ij}\}} \quad (3)$$

$$q'_x = \frac{\#\{d'_{XY} \leq b\}}{\#\{d'_{XY}\}} \quad (4)$$

where d_{ij} are distances within species, d_{XY} are distances among species for an *entire* genus of concern, and d'_{XY} are combined interspecific distances for a target species and its closest neighbouring species. The notation $\#$ reflects a count. Quantities a , a' , and b correspond to $\min(d_{XY})$, $\min(d'_{XY})$, and $\max(d_{ij})$, the minimum interspecific distance, the minimum combined interspecific distance, and the maximum intraspecific distance, respectively (**Figure 1**).




Figure 1: Modified depiction from Meyer and Paulay (2005) and Phillips et al. (2024) of the overlap/separation of intraspecific and interspecific distances (δ) for calculation of the DNA barcode gap metrics (p_x and q_x) for a hypothetical species x . The minimum interspecific distance is denoted by a and the maximum intraspecific distance is indicated by b . The quantity $f(\delta)$ is akin to a kernel density estimate of the probability density function of distances. A similar visualization can be displayed for p'_x and q'_x based on intraspecific and combined interspecific distances within the interval $[a', b]$.

Hence, Equations (1)-(4) are simply empirical partial means of distances falling at and below, or at and exceeding, given distribution thresholds. Notice further that a/a' , and b are also the first and n th order statistics, $X_{(1)}$ and $X_{(n)}$, respectively, with $a/a' < b$, which have been pointed out by Phillips et al. (2022) as important for developing a statistical theory to test for the existence of the DNA barcode gap. However, the underlying sampling distributions of $a - b$ and $a' - b$ are not obvious, thereby rendering a formal hypothesis test cumbersome to derive since said distributions are not expected to be symmetric. Equations (1)-(4) can also be expressed in terms of empirical cumulative distribution functions (ECDFs) (see next section). A clear connection between the ECDF and the nonparametric MLE was noted by Efron and Tibshirani (1993) as well as others. Distances form a continuous distribution and are easily computed from a model of DNA sequence evolution, such as uncorrected or corrected p-distances (Jukes and Cantor, 1969; Kimura, 1980) using, for example, the `dist.dna()` function available in the `ape` R package (Paradis et al., 2004). However, calculated values are *not* independent and identically distributed (IID) because estimated standard errors (SEs)

will depend on both the number of species sampled within the genus under study, as well as the number of specimens sampled within a target species. To tease this out, Phillips et al. (2024) suggests plotting estimator values against their estimated SEs, along with a simple random downsampling scheme. In the case of two species comprising a focal genus, one well sampled and the other poorly sampled, values of the metrics close to zero for the sufficiently sampled species will likely possess larger SEs following downsizing to match the number of poorly sampled specimens (Phillips et al., 2024).

The approach of Phillips et al. (2024) differs markedly from the traditional definition of the DNA barcoding gap laid out by Meyer and Paulay (2005) and Meier et al. (2008) in that the proposed metrics incorporate interspecific distances which *include* the target species of interest. Furthermore, if a focal species is found to have multiple nearest neighbours, then the species possessing the smallest average distance is used (Phillips et al., 2024). These schemes more accurately account for species' coalescence histories inferred from contemporaneous samples of DNA sequences leading to instances of barcode sequence sharing, such as interspecific hybridization/introgression events (Phillips et al., 2024). For example, within Equations (3) and (4), the degree of distance distribution overlap/separation between a target taxon and its nearest neighbouring species, gauged from magnitudes of p'_x and q'_x , is directly proportional to the amount of time in which the two lineages diverged from the MRCA (Phillips et al., 2024). More specifically, $a \propto t_{XY, \min}$ and $b \propto t_{ij, \max}$, the minimum and maximum (unobserved) divergence time (in N_e generations), respectively, for two different individuals (i, j) and species (X, Y) within a well-inventoried genus (Phillips et al., 2024). Figure 1 in Phillips et al. (2024) depicts a phylogeny for two hypothetical species, A and B , showing paraphyly based on the sequencing of two mitochondrial loci. The (unobserved) time the two species formed, t_{AB} , is intermediate between $t_{XY, \min}$ and $t_{ij, \max}$, but theoretically, its maximum ($t_{AB, \max}$) may occur more distantly in the past closer to the root comprising the MRCA for all sampled taxa (Phillips et al., 2024). From a coalescence perspective, p_x is the probability two lineages within species x coalesce after $t_{XY, \min}$, whereas q_x represents

the probability a lineage within species x coalesces with a lineage in species k (within a sample) after $t_{XY, \min}$, but before $t_{ij, \max}$. Thus, the quantities can be used as a criterion to assess the failure of DNA barcoding in recently radiated taxonomic groups, among other plausible biological explanations mentioned earlier. It should be noted that there is a loss of information in computing the values of Phillips et al.'s (2024) metrics since two different sets of (p_x, q_x) and (p'_x, q'_x) values cause equal but distinct underlying distance distributions of $f_{\text{intra}}(\delta)$ and $f_{\text{inter}}(\delta)$ (**Figure 1**). Note, distances are constrained to the interval $[0, 1]$, whereas the metrics are defined only on the interval $[a/a', b]$. Values of the estimators obtained from Equations (1)-(4) close to or equal to zero give evidence for separation between intraspecific and interspecific distance distributions; that is, values suggest the presence of a DNA barcode gap for a target species. Conversely, values near or equal to one give evidence for distribution overlap; that is, values likely indicate the absence of a DNA barcode gap.

2.2 The Model

Before delving into the derivation of the proposed DNA barcode gap metrics, review of some fundamental statistical theory is necessary.

For a given random variable X from a given distribution F , its cumulative distribution function (CDF) is defined by

$$F_X(t) = \mathbb{P}(X \leq t) = 1 - \mathbb{P}(X > t). \quad (5)$$

Rearranging Equation (5) gives

$$\mathbb{P}(X > t) = 1 - F_X(t) = 1 - \mathbb{P}(X \leq t), \quad (6)$$

from which it follows that

$$\mathbb{P}(X \geq t) = 1 - F_X(t) + \mathbb{P}(X = t). \quad (7)$$

The CDF has many desirable statistical properties, such as boundedness, continuity, and monotonicity. Consult any graduate-level text on mathematical statistics for reference.

Equations (1)-(4) can thus be written in terms of ECDFs via Equations (5) and (7) as follows, since the true underlying F s, are unknown *a priori*, and therefore must be estimated using available sequence distance data:

$$\begin{aligned} p_x &= \mathbb{P}(d_{ij} \geq a) \\ &= 1 - \hat{F}_{d_{ij}}(a) + \mathbb{P}(d_{ij} = a) \end{aligned} \quad (8)$$

$$\begin{aligned} q_x &= \mathbb{P}(d_{XY} \leq b) \\ &= \hat{F}_{d_{XY}}(b) \end{aligned} \quad (9)$$

$$\begin{aligned} p'_x &= \mathbb{P}(d_{ij} \geq a') \\ &= 1 - \hat{F}_{d_{ij}}(a') + \mathbb{P}(d_{ij} = a') \end{aligned} \quad (10)$$

$$\begin{aligned} q'_x &= \mathbb{P}(d'_{XY} \leq b) \\ &= \hat{F}_{d'_{XY}}(b). \end{aligned} \quad (11)$$

Given n increasing-ordered data points, the (discrete) ECDF, $\hat{F}_n(t) = \frac{1}{n} \sum_{i=1}^n \mathbb{1}_{[x_i \leq t]}$, comprises an increasing step function having jump discontinuities of size $\frac{1}{n}$ at each sample observation (x_i) , excluding ties (or steps of weight $\frac{i}{n}$ with duplicate observations), where $\mathbb{1}(x)$ is the indicator function. Note that $1 - \hat{F}_n(t)$ is the complementary ECDF (CECDF), which is a declining function. Further, notice that $\mathbb{P}(X = t) \neq 0$. The ECDF/CECDF has good asymptotic behaviour in the limit of sample size. Equations (8)-(11) clearly demonstrate the asymmetric directionality of the proposed metrics. Plots of the ECDFs/CECDFs for

nonzero intraspecific, interspecific, and combined interspecific distances, superimposed with corresponding values of a/a' and b , provide a new, useful way to visualize the DNA barcode gap. Firstly, there is no need to tune the output (*e.g.*, number of bins and bin width in the case of histograms). Secondly, probabilities reflecting areas under the distance curves depicted in **Figure 1** can be directly read off the ECDF/CECDF curves. Values of the metrics can also be determined from direct inspection of the ECDF/CECDF plot by determining the corresponding y -value (cumulative probability) for a given value of x (distance). Furthermore, calculation of the DNA barcode gap estimators is convenient as they implicitly account for total distribution area (including overlap).

A major criticism of large sample (frequentist) theory is that it relies on asymptotic properties of the MLE (whose population parameter is assumed to be a fixed but unknown quantity), such as estimator normality and consistency as the sample size approaches infinity. This problem is especially pronounced in the case of binomial proportions (Newcombe, 1998). The estimated Wald SE of the sample proportion, is given by $\widehat{SE}[\hat{p}] = \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$, where $\hat{p} = \frac{Y}{n}$ is the MLE, Y is the total number of successes ($Y = \sum_{i=1}^n y_i$), and n is the total number of trials (*i.e.*, sample size). However, the above formula for the standard error is problematic for several reasons. First, it is a Normal approximation which makes use of the central limit theorem (CLT); thus, large sample sizes are required for reliable estimation. When few observations are available, SEs will be large and inaccurate, leading to low statistical power to detect a true DNA barcode gap when one actually exists. Further, resulting interval estimates could span values less than zero or greater than one, or have zero width, which is practically meaningless. Second, when proportions are exactly equal to zero or one, resulting SEs will be exactly zero, rendering $\widehat{SE}[\hat{p}]$ given above completely useless. In the context of the proposed DNA barcode gap metrics, values obtained at the boundaries of their support are often encountered. Therefore, reliable calculation of SEs is not feasible. Given the importance of sufficient sampling of species genetic diversity for DNA barcoding initiatives, a different statistical estimation approach is necessary.

Bayesian inference offers a natural path forward in this regard since it allows for straightforward specification of prior beliefs concerning unknown model parameters and permits the seamless propagation of uncertainty, when data are lacking and sample sizes are small, through integration with the likelihood function associated with true generating processes. The posterior distribution ($\pi(\theta|Y)$) is given by Bayes' theorem up to a proportionality $\pi(\theta|Y) \propto \pi(Y|\theta)\pi(\theta)$, where θ are unobserved parameters, Y are known data, $\pi(Y|\theta)$ is the likelihood, and $\pi(\theta)$ is the prior. As a consequence, because parameters are treated as random variables, Bayesian models are much more flexible and generally more easily interpretable compared to frequentist approaches. Under the Bayesian paradigm, entire posterior distributions, along with their summaries (*e.g.*, CrIs) are outputted, rather than just long run behaviour reflected in sampling distributions, p-values, and CIs as in the frequentist case, thus allowing direct probability statements to be made.

Essentially, from a statistical perspective, the goal herein is to nonparametrically estimate probabilities corresponding to extreme tail quantiles for positive highly skewed distributions on the unit interval (or any closed subinterval thereof). Here, it is sought to numerically approximate the extent of proportional overlap/separation of intraspecific and interspecific distance distributions within the subinterval $[a/a', b]$. This is a challenging computational problem within the current study as detailed in subsequent sections. The usual approach employs kernel density estimation (KDE), along with numerical or Monte Carlo integration of complicated probability distribution functions (PDFs), and invocation of extreme value theory (EVT); however, this requires careful selection of the bandwidth parameter, which would need to be identical for both distributions, among other considerations. This becomes problematic when fitting finite mixture models using Markov Chain Monte Carlo (MCMC) where nonidentifiability is rampant due to the occurrence of label switching, making clusters hard to discern across sampling iterations, and model selection challenging due to lack of exchangeability among observations (Stephens, 2000). For DNA barcode gap estimation, this would correspond to a two-component mixture (one for intraspecific distance comparisons,

and the other for interspecific comparisons), with one or more distribution modes or curve intersection points between components, and the presence of zero distance inflation. This makes parameter estimation difficult using methods like the Expectation-Maximization (EM) algorithm (Dempster et al., 1977) as the algorithm may become stuck in suboptimal regions of the parameter search space and prematurely converge to local optima. This seems wholly plausible, especially given that the genus-level distribution of genetic distances is a mixture of closely-related (*i.e.*, nearest neighbour) and more distant species. Examining **Figure 1**, depending on divergence times and other elements, intraspecific and interspecific distributions would be expected to comprise greater density in the upper (right) tails (which make up more distantly-related taxa relative to a target species) compared to the lower (left) tails (which contain only the closest conspecifics) since there are more combinations of distance comparisons possible. A key problem affecting reliable inference of the DNA barcode gap metrics is undersampling of genus-level and/or interspecific variation. An immediate, but complex solution appears to be a move toward next-generation barcoding incorporating multiple loci. In theory, this would facilitate the detection of missing operational taxonomic units (OTUs), both extant and extinct, when attempting to better characterize a genus. Phillips et al.'s (2024) could be informative in this regard for studies of human phylogeography and ancestry composition where historical admixture between sympatric populations was commonplace (*e.g.*, modern human/Denisova interbreeding). Here, for simplicity, an alternate route is taken to avoid these obstacles.

Counts, y , of overlapping distances (as expressed in the numerator of Equations (1)-(4)) are treated as binomially distributed with expectation $\mathbb{E}[Y|\theta] = k\theta$, and variance $\mathbb{V}[Y|\theta] = k\theta(1 - \theta)$, where $k = \{N, C\}$ are total count vectors of intraspecific and combined interspecific distances, respectively, for a target species along with its nearest neighbour species, and $k = M$ is a total count vector for all interspecific species comparisons. This follows from the fact that the ECDF/CECDF is binomially distributed. The quantity thus being estimated is the parameter vector $\underline{\theta} = \{p_x, q_x, p'_x, q'_x\}$. The metrics encompassing $\underline{\theta}$ are

396 presumed to follow a $\text{Beta}(\alpha, \beta)$ distribution, with real shape parameters α and β , which
 397 is a natural choice of prior on probabilities. The beta distribution has a prior mean of
 398 $\mathbb{E}[\theta|\alpha, \beta] = \frac{\alpha}{\alpha+\beta}$ and a prior variance equal to $\mathbb{V}[\theta|\alpha, \beta] = \frac{\alpha\beta}{(\alpha+\beta)^2(\alpha+\beta+1)}$. In the case where
 399 $\alpha = \beta$, all generated $\text{Beta}(\alpha, \beta)$ distributions will possess the same prior expectation, whereas
 400 the prior variance will shrink as both α and β increase. Such a scheme is quite convenient
 401 since the beta distribution is conjugate to the binomial distribution. Thus, the posterior
 402 distribution is also beta distributed, specifically, $\text{Beta}(\alpha + Y, \beta + n - Y)$, having expectation
 403 $\mathbb{E}[\theta|Y, \alpha, \beta] = \frac{\alpha+Y}{\alpha+\beta+n}$ and variance $\mathbb{V}[\theta|Y, \alpha, \beta] = \frac{(\alpha+Y)(\beta+n-Y)}{(\alpha+\beta+n)^2(\alpha+\beta+n+1)}$. In the context of DNA
 404 barcoding, it is important that the DNA barcode gap metrics effectively differentiate between
 405 extremes of no overlap/complete separation and full overlap/no separation,
 406 corresponding to values of the metrics equal to 0 and 1 (equivalent to total distance counts
 407 of 0 and n), respectively. These extremes yield a posterior expectation of
 408 $\mathbb{E}[\theta|Y = 0, \alpha, \beta] = \frac{\alpha}{\alpha+\beta+n}$ and a posterior variance of $\mathbb{V}[\theta|Y = 0, \alpha, \beta] = \frac{\alpha(\beta+n)}{(\alpha+\beta+n)^2(\alpha+\beta+n+1)}$
 409 and $\mathbb{E}[\theta|Y = n, \alpha, \beta] = \frac{\alpha+n}{\alpha+\beta+n}$ and $\mathbb{V}[\theta|Y = n, \alpha, \beta] = \frac{(\alpha+n)\beta}{(\alpha+\beta+n)^2(\alpha+\beta+n+1)}$. Note, the posterior
 410 variances are equivalent at these thresholds for all $\alpha = \beta$.

411 Parameters were given an uninformative $\text{Beta}(1, 1)$ prior, which is identical to a standard
 412 uniform ($\text{Uniform}(0, 1)$) prior since it places equal probability on all parameter values within
 413 its support. This distribution has an expected value of $\mu = \frac{1}{2} = 0.500$ and a variance of
 414 $\sigma^2 = \frac{1}{12} \approx 0.083$. Further, the posterior is $\text{Beta}(Y+1, n-Y+1)$, from which various moments
 415 such as the expected value $\mathbb{E}[\theta|Y] = \frac{Y+1}{n+2}$ and variance $\mathbb{V}[\theta|Y] = \frac{(Y+1)(n-Y+1)}{(n+2)^2(n+3)}$, and other
 416 quantities, can be easily calculated. Clearly, $\mathbb{E}[\theta|Y = 0] = \frac{1}{n+2}$ and $\mathbb{V}[\theta|Y = 0] = \frac{n+1}{(n+2)^2(n+3)}$,
 417 and $\mathbb{E}[\theta|Y = n] = \frac{n+1}{n+2}$ and $\mathbb{V}[\theta|Y = n] = \frac{n+1}{(n+2)^2(n+3)}$. For example, with a single distance
 418 ($n = 1$) for a given species (wherein two specimens are compared), and where the distance
 419 is not overlapping, on average, θ will be equal to $\mu = \frac{1}{3} = 0.333$, with a variance of
 420 $\sigma^2 = \frac{2}{36} = \frac{1}{18} \approx 0.056$. In contrast, with $n = 100$ individuals and no overlap, the mean reduces
 421 to $\mu = \frac{1}{102} \approx 0.001$, with a variance of $\sigma^2 = \frac{101}{1071612} = 9.425 \times 10^{-5}$. When $n = 1$ and there is
 422 overlap, the mean is $\mu = \frac{2}{3} = 0.667$, having a variance of $\sigma^2 \approx 0.0556$. For $n = 100$ where all

distances overlap, θ has expectation $\mu = \frac{101}{102} \approx 0.990$ and variance $\sigma^2 \approx 9.425 \times 10^{-5}$. These edge case examples demonstrate that the Beta(1, 1) distribution serves as a good first-pass prior for DNA barcode gap estimation using Phillips et al.'s (2024) metrics, though tighter bounds can likely be achieved. In general however, when possible, it is always advisable to incorporate prior information, even if only weak, rather than simply imposing complete ignorance in the form of a flat prior distribution. In the case of unimodal distributions, the (estimated) posterior mean often possesses the property that it readily decomposes into a convex linear combination, in the form of a weighted sum, of the (estimated) prior mean and the MLE. That is $\hat{\mu}_{\text{posterior}} = w\hat{\mu}_{\text{prior}} + (1 - w)\hat{\mu}_{\text{MLE}}$, where for the beta distribution, $w = \frac{\alpha + \beta}{\alpha + \beta + n}$. Therefore, with sufficient data, $w \rightarrow 0$ as $n \rightarrow \infty$, regardless of the values of α and β , and the choice of prior distribution becomes less important since the posterior will be dominated by the likelihood. For the Beta(1, 1), $w = \frac{2}{2+n}$, with $n = 2$ giving $w = \frac{1}{2} = 0.500$; that is, the posterior is the arithmetic average of the prior and the likelihood. For $n = 100$, the prior weight is $w = \frac{2}{102} \approx 0.0196$. The full Bayesian model for species x is thus given by

$$\begin{aligned}
y_{\text{lwr}} &\sim \text{Binomial}(N, p_{\text{lwr}}) \\
y_{\text{upr}} &\sim \text{Binomial}(M, p_{\text{upr}}) \\
y'_{\text{lwr}} &\sim \text{Binomial}(N, p'_{\text{lwr}}) \\
y'_{\text{upr}} &\sim \text{Binomial}(C, p'_{\text{upr}})
\end{aligned} \tag{12}$$

$$p_{\text{lwr}}, p_{\text{upr}}, p'_{\text{lwr}}, p'_{\text{upr}} \sim \text{Beta}(1, 1).$$

Note that p_x , q_x , p'_x , and q'_x in Equations (1)-(4) are denoted p_{lwr} , p_{upr} , p'_{lwr} , and q'_{upr} , respectively within Equation (12) for easy distinction between MLEs and Bayesian posterior estimates. The above statistical theory and derivations lay a good foundation for the remainder of this paper.

The proposed model is inherently vectorized to allow processing of multiple species

datasets simultaneously. Model fitting was achieved using the Stan probabilistic
 programming language (Carpenter et al., 2017) framework for Hamiltonian Monte Carlo
 (HMC) (Neal, 2011) via the No-U-Turn Sampler (NUTS) sampling algorithm (Hoffman and
 Gelman, 2014) through the `rstan` R package (version 2.32.6) (Stan Development Team, 2023)
 in R (version 4.4.1) (R Core Team, 2024). Four Markov chains were run for 2000 iterations
 each in parallel across four cores with random parameter initializations. Within each chain, a
 total of 1000 samples was discarded as warmup (*i.e.*, burnin) to reduce dependence on starting
 conditions and to ensure posterior samples are reflective of the equilibrium distribution.
 Further, 1000 post-warmup draws were utilized per chain during the sampling phase. Because
 HMC/NUTS results in dependent samples that are minimally autocorrelated, chain thinning
 is not required. Each of these tuning parameters reflect default MCMC settings in Stan
 to control both bias and variance, respectively, in the resulting draws. All analyses in the
 present work were carried out on a 2023 Apple MacBook Pro with M2 chip and 16 GB RAM
 running macOS Ventura 13.2. A random seed was set to ensure reproducibility of model
 results. Outputted estimates were rounded to three decimal places of precision. Posterior
 distributions were visualized as KDE plots using the `ggplot()` function within the `ggplot2`
 R package (version 3.5.1) (Wickham, 2016) with the default Gaussian kernel and optimal
 smoothness selection. To successfully run the Stan program, end users must have installed
 an appropriate compiler (such as GCC or Clang) which is compatible with their operating
 system, such as macOS.

Convergence was assessed both visually and quantitatively as follows: (1) through
 examining parameter traceplots, which depict the trajectory of accepted MCMC draws as
 a function of the number of iterations, (2) through monitoring the Gelman-Rubin potential
 scale reduction factor statistic ($\hat{R}$) (Gelman and Rubin, 1992; Vehtari et al., 2021), which
 measures the concordance of within-chain *versus* between-chain variance, and (3) through
 calculating the effective sample size (ESS) for each parameter, which quantifies the number
 of independent samples generated Markov chains are equivalent to. Mixing of chains was

deemed sufficient when traceplots looked like “fuzzy caterpillars”, $\hat{R} < 1.01$, and effective sample sizes were reasonably large (Gelman et al., 2020). After sampling, a number of summary quantities were reported, including posterior means, posterior SDs, and posterior quantiles from which 95% CrIs could be computed to make probabilistic inferences concerning true population parameters. To validate the overall correctness of the proposed statistical model given by Equation (12), as a means of comparison, posterior predictive checks (PPCs) were also employed to generate binomial random variates in the form of counts from the posterior predictive distribution; that is $\gamma = \{Np_x, Mq_x, Np'_x, Cq'_x\}$ to verify that the model adequately captures relevant features of the observed data. The proposed Bayesian model outlined here has a straightforward interpretation (**Table 1**).

3 Results and Discussion

The *Agabus* CYTB dataset analyzed by Phillips et al. (2024) is revisited herein for the species *A. bipustulatus* and *A. nevadensis*, since these taxa were the sole representatives for this locus, with the most and the least specimen records, respectively ($N = 701$ and $N = 2$) across all three assessed molecular markers. Briefly, using the R package **MACER** (Young et al., 2021), DNA sequences were downloaded from GenBank and BOLD and processed using a stringent workflow to obtain a 343-bp high-quality FASTA alignment representing 46 unique haplotypes. The resulting alignment contained no stop codons, ambiguous nucleotides, or other problematic artifacts that bias genetic diversity estimates. Further, both species- and genus-level distance outliers were excluded. Genetic distances were calculated using uncorrected p-distances. *A. bipustulatus* comprised 45 total haplotypes, whereas *A. nevadensis* possessed only one haplotype. Note, DNA barcode gap estimation is only possible for species having at least two specimen records because only one distance can be calculated. This dataset is a prime illustrative example highlighting the issue of inadequate taxon sampling, which arises frequently in large-scale phylogenetic and phylogeographic

Table 1: Interpretation of the DNA barcode gap estimators within $[a/a', b]$

Parameter	Explanation
p_x/p_{lwr}	When p_x/p_{lwr} is close to 0 (1), it suggests that the probability of intraspecific (interspecific) distances being larger (smaller) than interspecific (intraspecific) distances is low (high) on average, while the probability of interspecific (intraspecific) distances being larger (smaller) than intraspecific (interspecific) distances is high (low) on average; that is, there is (no) evidence for a DNA barcode gap.
q_x/p_{upr}	When q_x/p_{upr} is close to 0 (1), it suggests that the probability of interspecific (intraspecific) distances being larger (smaller) than intraspecific (interspecific) distances is high (low) on average, while the probability of intraspecific (interspecific) distances being larger (smaller) than interspecific (intraspecific) distances is low (high) on average; that is, there is (no) evidence for a DNA barcode gap.
p'_x/p'_{lwr}	When p'_x/p'_{lwr} is close to 0 (1), it suggests that the probability of intraspecific (combined interspecific distances for a target species and its nearest neighbour species) distances being larger (smaller) than combined interspecific distances for a target species and its nearest neighbour species (intraspecific distances) is low (high) on average, while the probability of combined interspecific distances for a target species and its nearest neighbour species (intraspecific distances) being larger (smaller) than intraspecific distances (combined interspecific distances for a target species and its nearest neighbour species) is high (low) on average; that is, there is (no) evidence for a DNA barcode gap.
q'_x/p'_{upr}	When q'_x/p'_{upr} is close to 0 (1), it suggests that the probability of combined interspecific distances for a target species and its nearest neighbour species (intraspecific distances) being larger (smaller) than intraspecific distances (combined interspecific distances for a target species and its nearest neighbour species) is high (low) on average, while the probability of intraspecific distances (combined interspecific distances for a target species and its nearest neighbour species) being larger (smaller) than combined interspecific distances for a target species and its nearest neighbour species (intraspecific distances) is low (high) on average; that is, there is (no) evidence for a DNA barcode gap.

studies, in several respects. First, from a statistical viewpoint, sample sizes reflect extremes in reliable parameter estimation. Second, from a DNA barcoding perspective, *Agabus* currently comprises about 200 extant Linnean species according to the Global Biodiversity Information Facility (GBIF) (<https://www.gbif.org>); yet, due to the level of convenience sampling inherent in taxonomic collection efforts for this genus, adequate representation of species and genetic diversity is far from complete. Indeed, herein, species coverage for CYTB is only around 1.000%. Not surprisingly, CYTB performed poorly as a DNA barcode locus for *Agabus* despite having the highest specimen representativeness for a given species ($N = 701$, *A. bipustulatus*) and displaying good separation of genetic sequences among all three mitochondrial genes examined in Phillips et al.'s (2024) study; yet no barcode gap was observed since all metrics were found to be equal to 1.000.

ECDF/CECDF plots were generated using the base R function `ecdf()`, along with `ggplot()`, and are displayed in **Figure 2**. Interpretation is simple: in the case of *A. bipustulatus*, p_x , all distances are greater than $a = 0$, with 36.900% of distances equal to a . Conversely, both q_x and q'_x equal one because all interspecific distances are less than $b \approx 0.026$; for *A. nevadensis*, both metrics are approximately equal to 0.010 since only about 1% of distances fall below $b = 0$.




Figure 2: ECDF/CECDF plots of the DNA barcode gap metrics for *A. bipustulatus* and *A. nevadensis* with a/a' and b clearly indicated as dashed vertical lines. When all genetic distances are equal to zero, no step function is shown.

511 MCMC parameter traceplots showed rapid mixing of chains to the stationary distribution
 512 (**Figure 3**). Further, all $\hat{R}$ and ESS values (not shown) were close to their recommended
 513 cutoffs of one and thousands of samples, respectively, indicating chains are both well-mixed
 514 and have converged to the posterior distribution. Bayesian posterior estimates were reported
 515 alongside frequentist MLEs, in addition to SEs, posterior SDs, 95% CIs, and 95% CrIs
 516 (**Table 2**).




Figure 3: MCMC parameter traceplots applied to *A. bipustulatus* ([1]; $N = 701$) and *A. nevadensis* ([2]; $N = 2$) for CYTB across 1000 post-warmup iterations.

Table 2: Nonparametric frequentist and Bayesian estimates of distance distribution overlap/separation for the DNA barcode gap coalescent model parameters applied to *A. bipustulatus* ($N = 701$) and *A. nevadensis* ($N = 2$) for CYTB, including 95% CIs and CrIs. CrIs are based on 4000 posterior draws. All parameter estimates are reported to three decimal places of precision.

Species	Parameter	MLE (SE, 95% CI)	Bayes Est. (SD; 95% CrI)
<i>A. bipustulatus</i>	p_x/p_{lwr}	1.000 (0.000; 1.000-1.000)	1.000 (0.000; 1.000-1.000)
<i>A. bipustulatus</i>	q_x/p_{upr}	1.000 (0.000; 1.000-1.000)	1.000 (0.000; 0.999-1.000)
<i>A. bipustulatus</i>	p'_x/p'_{lwr}	1.000 (0.000; 1.000-1.000)	1.000 (0.000; 1.000-1.000)
<i>A. bipustulatus</i>	q'_x/p'_{upr}	1.000 (0.000; 1.000-1.000)	1.000 (0.000; 0.999-1.000)
<i>A. nevadensis</i>	p_x/p_{lwr}	1.000 (0.000; 1.000-1.000)	0.835 (0.144; 0.470-0.996)
<i>A. nevadensis</i>	q_x/p_{upr}	0.010 (0.002; 0.006-0.014)	0.010 (0.002; 0.007-0.014)
<i>A. nevadensis</i>	p'_x/p'_{lwr}	1.000 (0.000; 1.000-1.000)	0.834 (0.138; 0.481-0.994)
<i>A. nevadensis</i>	q'_x/p'_{upr}	0.010 (0.070; -0.128-0.148)	0.010 (0.002; 0.007-0.014)

CIs were calculated using the usual large sample $(1 - \alpha)100\%$ -level interval estimate given by $\hat{p} \pm z_{1-\frac{\alpha}{2}} \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$, where $z_{1-\frac{\alpha}{2}} = 1.960$ is the critical value for 95% confidence (*i.e.*, the 97.5th percent quantile from the standard Normal distribution), and α is the stated significance level (here, 5%). Given a $(1 - \alpha)100\%$ CI, with repeated sampling, on average $(1 - \alpha)100\%$ of constructed intervals will contain the true parameter of interest, assuming a correctly-specified statistical model; on the other hand, any given CI will either capture or exclude the true parameter with 100% certainty. This in stark contrast to a CrI, where the true parameter is contained within said interval with $(1 - \alpha)100\%$ probability, contingent on the data generating model being correct. Note, by default Stan computes equal-tailed (central) CrIs such that there is equal area situated in the left and right tails of the posterior distribution. For a 95% CrI, this corresponds to the 2.5th and 97.5th percent quantiles. However, constructed intervals are usually only valid for symmetric or nearly symmetric distributions. Given the bounded nature of the DNA barcode gap metrics, whose posterior distributions, as expected, show considerable skewness, an alternative approach to reporting CrIs, such as Highest Posterior Density (HPD) intervals (Chen and Shao, 1999) or shortest probability intervals (SPIn) (Liu et al., 2015) is warranted. Although the data for *Agabus* was not found to be consistent with a value of zero for the metrics (suggesting the existence of a DNA barcode gap) for either species, this is clearly a case that must be differentiated from

one where some or all of the estimators attain a value of one (indicating no DNA barcode gap likely exists). As such asymmetric intervals generally attain greater statistical efficiency (in the form of smaller Mean Squared Error (MSE) or variance) and higher coverage probabilities than more standard interval estimates, careful in-depth comparison is left for future work.

Findings based on nonparametric MLEs and Bayesian posterior means were quite comparable with one another and show evidence of complete overlap in intraspecific, interspecific, and combined interspecific distances for *A. bipustulatus* in both the $p/q/p_{\text{LWR}}/p_{\text{UPR}}$ and $p'/q'/p'_{\text{LWR}}/p'_{\text{UPR}}$ directions since the metrics and their intervals attain magnitudes very close or equal to to one, with minimal noise (**Table 2**). As a result, this likely indicates that no DNA barcode gap is present for this species. Such findings are strongly reinforced by the very tight clustering of posterior draws (**Figure 4**) and associated interval estimates owing to the large number of specimens sampled for this species.




Figure 4: Scatterplots (black solid points) and distributions (black solid lines) depicting the DNA barcode gap metrics for *A. bipustulatus* ($N = 701$) and *A. nevadensis* ($N = 2$) across CYTB based on 4000 Bayesian posterior draws. MLEs and posterior means are displayed as coloured (red/blue) solid and dashed lines for the metrics, respectively.

On the other hand, the situation for *A. nevadensis* is more nuanced, as posterior values are further spread out (**Table 2** and **Figure 4**), suggesting less overall certainty in true parameter values given the low specimen sampling coverage for this taxon. Of note, calculated SEs and SDs are large, with the 95% CIs and 95% CrIs being quite wide in most cases, consistent with much uncertainty regarding the computed frequentist and Bayesian posterior means of the DNA barcode gap metrics. For instance, the Bayesian analysis suggests that the data are consistent with both p_{lwr} and p'_{lwr} ranging from approximately 0.500-1.000 due to the large SD of about 0.140. The posterior means associated with these estimates themselves are also far from one. The CrI for p_{upr} and p'_{upr} also spans an order of magnitude, with the lower tail estimate of 0.007 being very close to zero. Further, regarding the frequentist analysis for the same species, the 95% CI for q_x is quite wide, reflecting considerable uncertainty in its true parameter value. Similarly, that for q'_x extends to negative values at the left endpoint, due to the corresponding SE of 0.070 being too high as a result of the extremely low sample size of $n = 2$ individuals sampled (**Table 2**). Since the 95% CI for q'_x truncated at the lower endpoint includes the value of zero, the null hypothesis for the presence of a DNA barcode gap cannot be rejected. Despite this, it is worth noting that truncation is not standard statistical practice and will likely lead to an interval with less than 95% nominal coverage having high relative error. In such cases, more appropriate confidence interval methods like the Wilson score interval, the exact (Clopper-Pearson) interval, or the Agresti-Coull interval should be employed (Newcombe, 1998; Agresti and Coull, 1998). KDEs for *A. bipustulatus* are strongly left (negatively) skewed (**Figure 4**), whereas those for *A. nevadensis* exhibit more symmetry, especially for p_{upr} and p'_{upr} (**Figure 4**). These differences are likely due to the stark contrast in sample sizes for the two examined species. Nevertheless, simulated counts of overlapping specimen records from the posterior predictive distribution (**Table 3**) were found to be very close to observed counts for both species, indicating that the proposed model adequately captures underlying variation.

Table 3: Posterior predictive checks of distance distribution overlap/separation for the DNA barcode gap coalescent model parameters applied to *A. bipustulatus* ($N = 701$) and *A. nevadensis* ($N = 2$) for CYTB. CrIs are based on 4000 posterior draws. All parameter estimates are reported to three decimal places of precision.

Species	Variable	Y	n	Bayes Est. (SD; 95% CrI)
<i>A. bipustulatus</i>	y_{lwr}	491401.000	491401.000	491400.018 (1.378; 491396.000-491401.000)
<i>A. bipustulatus</i>	y_{upr}	2804.000	2804.000	2803.019 (1.433; 2799.000-2804.000)
<i>A. bipustulatus</i>	y'_{jwr}	491401.000	491401.000	491400.008 (1.412; 491396.000-491401.000)
<i>A. bipustulatus</i>	y_{upr}	2804.000	2804.000	2802.992 (1.429; 2799.000-2804.000)
<i>A. nevadensis</i>	y_{lwr}	4.000	4.000	3.355 (0.888; 1.000-4.000)
<i>A. nevadensis</i>	y_{upr}	28.000	2804.000	29.151 (7.620; 16.000-45.000)
<i>A. nevadensis</i>	y'_{jwr}	4.000	4.000	3.325 (0.884; 1.000-4.000)
<i>A. nevadensis</i>	y_{upr}	28.000	2804.000	28.942 (7.409; 15.000-45.000)

Obtained results suggest that use of the Beta(1, 1) prior may not be appropriate given a low number of collected individuals for most taxa in DNA barcoding efforts. This suggests that further consideration of more informative beta priors is worthwhile. This is clear for *A. nevadensis*, where use of a stronger prior will likely eliminate the issue regarding accurate estimation of p_{lwr} and p'_{lwr} and stabilize their uncertainty.

4 Conclusion

Herein, the accuracy of the DNA barcode gap was analyzed from a rigorous statistical lens to expedite both the curation and growth of reference sequence libraries, ensuring they are populated with high quality, statistically defensible specimen records fit for purpose to address standing questions in ecology, evolutionary biology, management, and conservation. To accomplish this, recently proposed, easy to calculate, nonparametric MLEs were formally derived using ECDFs/CECDFs and applied to assess the extent of overlap/separation of distance distributions within and among two species of predatory water beetles in the genus *Agabus* sequenced at CYTB using a Bayesian binomial count model with conjugate beta priors. Findings highlight a high level of parameter uncertainty for *A. nevadensis*, whereas posterior estimates of the DNA barcode gap metrics for *A. bipustulatus* are much more certain. Based on these results, it is imperative that specimen sampling be prioritized to better reflect actual species boundaries. More generally, apart from the metrics being

employed to better highlighting the importance of within-species genetic diversity versus between-species divergence, it is expected that the approach developed herein will be of broad utility in applied fields, such as DNA-based detection of seafood fraud within global supply chains, and in the determination of species occupancy/detection probabilities at ecological sites of interest using active and passive environmental DNA (eDNA) methods such as metabarcoding.

Since the DNA barcode gap metrics often attain values very close to zero (suggesting no overlap and complete separation of distance distributions) and/or very near one (indicating no separation and complete overlap), in addition to more intermediate values, a noninformative $\text{Beta}(\frac{1}{2}, \frac{1}{2})$ prior may be more appropriate over complete ignorance imposed by a $\text{Beta}(1, 1)$ prior. The former distribution is U-shaped symmetric and places greater probability density at the extremes of the distribution due to its heavier tails, while still allowing for variability in parameter estimates within intermediate values along its domain. Note that this prior is Jeffreys' prior density (Jeffreys, 1946), which is proportional to the square root of the Fisher information $\mathcal{I}(\theta)$; that is, $\pi(\theta) \propto \theta^{-\frac{1}{2}}(1 - \theta)^{-\frac{1}{2}}$. Jeffreys' prior has several desirable statistical properties as a prior: that it is inversely proportional to the standard deviation of the binomial distribution, and most notably, that it is invariant to model reparameterization (Gelman et al., 2014). However, this prior can lead to divergent transitions, among other pathologies, imposed by complex geometry (*i.e.*, curvature) in the posterior space since many iterative stochastic MCMC sampling algorithms experience difficulties when exploring high density distribution regions. Thus, remedies to resolve them, such as lowering the step size of the HMC/NUTS sampler, should be attempted in future work, along with other approaches such as empirical Bayes estimation to approximate beta prior hyperparameters from observed data through the MLE or other methods of parameter estimation, such as the method of moments. Alternatively, hierarchical modelling could be employed to estimate separate distribution model hyperparameters for each species and/or compute distinct estimates for the directionality/comparison level of the DNA barcode gap metrics (*i.e.*, lower *vs.* upper,

618 non-prime *vs.* prime) separately within the genus under study. This would permit greater
619 flexibility through incorporating more fine-grained structure seen in the data; however, low
620 taxon sample sizes may preclude valid inferences to be reasonably ascertained due to the
621 large number additional parameters which would be introduced through the specification
622 of the hyperprior distributions. Methods outlined in Gelman et al. (2014, 2020), such as
623 dealing with non-exchangeability of observations and alternate model parameterizations like
624 the logit, may prove useful in this regard. Perhaps, the best course of action though is
625 the elicitation of priors from domain experts. Even though more work remains, it is clear
626 that both frequentist and Bayesian inference hold much promise for the future of molecular
627 biodiversity science.

Supplementary Information

None declared

Data Availability Statement

Raw data, R, and Stan code can be accessed via Dryad at:

<http://datadryad.org/stash/share/>

RZIfMixcEODe0RWP7eyXWQewSVbqEIA9UTrH3ZVKyn4.

A GitHub repository can be found at:

<https://github.com/jphill01/Bayesian-DNA-Barcode-Gap-Coalescent>.

Acknowledgements

We wish to recognise the valuable comments and discussions of Daniel (Dan) Gillis, Robert (Bob) Hanner, Robert (Rob) Young, and XXX anonymous reviewers.

We acknowledge that the University of Guelph resides on the ancestral lands of the Attawandaron people and the treaty lands and territory of the Mississaugas of the Credit. We recognize the significance of the Dish with One Spoon Covenant to this land and offer our respect to our Anishinaabe, Haudenosaunee and Métis neighbours as we strive to strengthen our relationships with them.

Funding

None declared.

Conflict of Interest

None declared.

Author Contributions

JDP wrote the manuscript, wrote R and Stan code, as well as analyzed and interpreted all model results.

References

Agresti, A. and B. A. Coull

1998. Approximate is better than ‘exact’ for interval estimation of binomial proportions. *The American Statistician*, 52(2):119–126.

Ahrens, D., F. Fujisawa, H.-J. Krammer, J. Eberle, S. Fabrizi, and A. Vogler

2016. Rarity and incomplete sampling in DNA-based species delimitation. *Systematic Biology*, 65(3):478–494.

Avise, J., J. Arnold, R. Ball, Jr., E. Bermingham, T. Lamb, J. Neigel, C. Reeb, and N. Saunders

1987. Intraspecific phylogeography: The mitochondrial DNA bridge between population genetics and systematics. *Annu. Rev. Ecol. Syst.*, 18:489–522.

Bartlett, S. and W. Davidson

1992. FINS (forensically informative nucleotide sequencing): A procedure for identifying the animal origin of biological specimens. *BioTechniques*, 12(3):408–411.

Bergsten, J., D. Bilton, T. Fujisawa, M. Elliott, M. Monaghan, M. Balke, L. Hendrich,

J. Geijer, J. Herrmann, G. Foster, I. Ribera, A. Nilsson, T. Barraclough, and A. Vogler

2012. The effect of geographical scale of sampling on DNA barcoding. *Systematic Biology*, 61(5):851–869.

Bouckaert, R., T. G. Vaughan, J. Barido-Sottani, S. Duchêne, M. Fourment,

A. Gavryushkina, J. Heled, G. Jones, D. Kühnert, N. De Maio, M. Matschiner, F. K.

Mendes, N. F. Müller, H. A. Ogilvie, L. du Plessis, A. Poppinga, A. Rambaut, D. Rasmussen,
I. Siveroni, M. A. Suchard, C.-H. Wu, D. Xie, C. Zhang, T. Stadler, and A. J. Drummond
2019. BEAST 2.5: An advanced software platform for Bayesian evolutionary analysis.
PLOS Computational Biology, 15(4):1–28.

Bucklin, A., D. Steinke, and L. Blanco-Bercial
2011. Dna barcoding of marine metazoa. *Annual Review of Marine Science*, 3(Volume 3,
2011):471–508.

Čandek, K. and M. Kuntner
2015. DNA barcoding gap: Reliable species identification over morphological and
geographical scales. *Molecular Ecology Resources*, 15(2):268–277.

Carpenter, B., A. Gelman, M. Hoffman, D. Lee, B. Goodrich, M. Betancourt, M. Brubaker,
J. Guo, P. Li, and A. Riddell
2017. Stan: A probabilistic programming language. *Journal of Statistical Software*, 76:1.

Carstens, B. C., T. A. Pelletier, N. M. Reid, and J. D. Satler
2013. How to fail at species delimitation. *Molecular Ecology*, 22(17):4369–4383.

Chen, M.-H. and Q.-M. Shao
1999. Monte Carlo estimation of Bayesian credible and HPD intervals. *Journal of
Computational and Graphical Statistics*, 8(1):69–92.

Collins, R. A. and R. H. Cruickshank
2013. The seven deadly sins of DNA barcoding. *Molecular Ecology Resources*,
13(6):969–975.

Collins, R. A. and R. H. Cruickshank
2014. Known knowns, known unknowns, unknown unknowns and unknown knowns in dna
barcoding: A comment on dowton et al. *Systematic Biology*, 63(6):1005–1009.

695 Dayrat, B.
696 2005. Towards integrative taxonomy. *Biological Journal of the Linnean Society*,
697 85(3):407–417.

698 de Queiroz, K.
699 2007. Species concepts and species delimitation. *Systematic Biology*, 56(6):879–886.

700 Dempster, A. P., N. M. Laird, and D. B. Rubin
701 1977. Maximum likelihood from incomplete data via the em algorithm. *Journal of the*
702 *Royal Statistical Society: Series B (Methodological)*, 39(1):1–22.

703 Downton, M., K. Meiklejohn, S. L. Cameron, and J. Wallman
704 2014. A preliminary framework for DNA barcoding, incorporating the multispecies
705 coalescent. *Systematic Biology*, 63(4):639–644.

706 Efron, B. and R. J. Tibshirani
707 1993. *An Introduction to the Bootstrap*. Chapman & Hall.

708 Flouri, T., X. Jiao, B. Rannala, and Z. Yang
709 2018. Species tree inference with BPP using genomic sequences and the multispecies
710 coalescent. *Molecular Biology and Evolution*, 35(10):2585–2593.

711 Fonseca, E. M., D. J. Duckett, and B. C. Carstens
712 2021. P2C2M.GMYC: An R package for assessing the utility of the Generalized Mixed
713 Yule Coalescent model. *Methods in Ecology and Evolution*, 12(3):487–493.

714 Fujisawa, T. and T. Barraclough
715 2013. Delimiting species using single-locus data and the generalized mixed yule coalescent
716 approach: A revised method and evaluation on simulated data sets. *Systematic Biology*,
717 62(5):707–724.

718 Fujita, M. K., A. D. Leaché, F. T. Burbrink, J. A. McGuire, and C. Moritz
719 2012. Coalescent-based species delimitation in an integrative taxonomy. *Trends in Ecology*
720 *& Evolution*, 27(9):480–488.

721 Gelman, A., J. Carlin, H. Stern, D. Duncan, A. Vehtari, and D. Rubin
722 2014. *Bayesian Data Analysis*, third edition. Chapman and Hall/CRC.

723 Gelman, A. and D. Rubin
724 1992. Inference from iterative simulation using multiple sequences. *Statistical Science*,
725 7(4):457–472.

726 Gelman, A., A. Vehtari, D. Simpson, C. Margossian, B. Carpenter, Y. Yao, L. Kennedy,
727 J. Gabry, P.-C. Bürkner, and M. Modrák
728 2020. Bayesian workflow.

729 Hart, M. and J. Sunday
730 2007. Things fall apart: biological species form unconnected parsimony networks. *Biology*
731 *Letters*, 3(5):509–512.

732 Hebert, P., A. Cywinska, S. Ball, and J. deWaard
733 2003a. Biological identifications through DNA barcodes. *Proceedings of the Royal Society*
734 *of London B: Biological Sciences*, 270(1512):313–321.

735 Hebert, P., S. Ratnasingham, and J. de Waard
736 2003b. Barcoding animal life: Cytochrome c oxidase subunit 1 divergences among
737 closely related species. *Proceedings of the Royal Society of London B: Biological Sciences*,
738 270(Suppl 1):S96–S99.

739 Hebert, P. D., M. Y. Stoeckle, T. S. Zemplak, and C. M. Francis
740 2004. Identification of birds through DNA barcodes. *PLoS Biol*, 2(10):e312.

741 Hickerson, M. J., C. P. Meyer, and C. Moritz
 742 2006. DNA barcoding will often fail to discover new animal species over broad parameter
 743 space. *Systematic Biology*, 55(5):729–739.

744 Ho, S. and S. Duchêne
 745 2014. Molecular-clock methods for estimating evolutionary rates and timescales. *Molecular*
 746 *Ecology*, 23(24):5947–5965.

747 Hoffman, M. and A. Gelman
 748 2014. The No-U-Turn Sampler: Adaptively setting path lengths in Hamiltonian Monte
 749 Carlo. *Journal of Machine Learning Research*, 15:1593–1623.

750 Hubert, N. and R. Hanner
 751 2015. DNA barcoding, species delineation and taxonomy: A historical perspective. *DNA*
 752 *Barcodes*, 3:44–58.

753 Hubert, N., J. Phillips, and R. Hanner
 754 2024. Delimiting species with single-locus dna sequences. In *DNA Barcoding*, R. DeSalle,
 755 ed., volume 2744 of *Methods in Molecular Biology*. New York, NY: Humana.

756 Jackson, N. D., B. C. Carstens, A. E. Morales, and B. C. O’Meara
 757 2017a. Species delimitation with gene flow. *Systematic Biology*, 66(5):799–812.

758 Jackson, N. D., A. E. Morales, B. C. Carstens, and B. C. O’Meara
 759 2017b. PHRAPL: Phylogeographic inference using approximate likelihoods. *Systematic*
 760 *Biology*, 66(6):1045–1053.

761 Jeffreys, H.
 762 1946. An invariant form for the prior probability in estimation problems. *Proceedings*
 763 *of the Royal Society of London. Series A, Mathematical and Physical Sciences*,
 764 186(1007):453–461.

765 Jukes, T. and C. Cantor
766 1969. Evolution of protein molecules. In *Mammalian Protein Metabolism*, H. N. Munro,
767 ed., Pp. 21–132. New York: Academic Press.

768 Kapli, P., S. Lutteropp, J. Zhang, K. Kobert, P. Pavlidis, A. Stamatakis, and T. Flouri
769 2017. Multi-rate poisson tree processes for single-locus species delimitation under maximum
770 likelihood and markov chain monte carlo. *Bioinformatics*, 33(11):1630–1638.

771 Kimura, M.
772 1980. A simple method for estimating evolutionary rates of base substitutions
773 through comparative studies of nucleotide sequences. *Journal of Molecular Evolution*,
774 16(1):111–120.

775 Kingman, J.
776 1982a. The coalescent. *Stochastic Processes and Their Applications*, 13:235–248.

777 Kingman, J.
778 1982b. On the genealogy of large populations. *Journal of Applied Probability*, 19(A):27–43.

779 Knowles, L. L.
780 2009. Statistical phylogeography. *Annual Review of Ecology, Evolution, and Systematics*,
781 40(Volume 40, 2009):593–612.

782 Knowles, L. L. and W. P. Maddison
783 2002. Statistical phylogeography. *Molecular Ecology*, 11(12):2623–2635.

784 Kullback, S. and R. Leibler
785 1951. On information and sufficiency. *Annals of Mathematical Statistics*, 22(1):79–86.

786 Leaché, A., M. Fujita, V. Minin, and R. Bouckaert
787 2014. Species delimitation using genome-wide snp data. *Systematic Biology*, 63(4):534–542.

Liu, Y., A. Gelman, and T. Zheng

2015. Simulation-efficient shortest probability intervals. *Statistical Computing*, 25:809–819.

Mather, N., S. Traves, and S. Ho

2019. A practical introduction to sequentially Markovian coalescent methods for estimating demographic history from genomic data. *Ecology and Evolution*, 10(1):579–589.

Mayr, E.

1942. *Systematics and the Origin of Species from the Viewpoint of a Zoologist*. New York: Columbia University Press.

Meier, R., G. Zhang, and F. Ali

2008. The use of mean instead of smallest interspecific distances exaggerates the size of the “barcoding gap” and leads to misidentification. *Systematic Biology*, 57(5):809–813.

Meyer, C. and G. Paulay

2005. DNA barcoding: Error rates based on comprehensive sampling. *PLOS Biology*, 3(12):e422.

Monaghan, M., R. Wild, M. Elliot, T. Fujisawa, M. Balke, D. Inward, D. Lees, R. Ranaivosolo, P. Eggleton, T. Barraclough, and A. Vogler

2009. Accelerated species inventory on madagascar using coalescent-based models of species delineation. *Systematic Biology*, 58:298–311.

Mutanen, M., S. M. Kivelä, R. A. Vos, C. Doorenweerd, S. Ratnasingham, A. Hausmann, P. Huemer, V. Dincă, E. J. van Nieukerken, C. Lopez-Vaamonde, R. Vila, L. Aarvik, T. Decaëns, K. A. Efetov, P. D. N. Hebert, A. Johnsen, O. Karsholt, M. Pentinsaari, R. Rougerie, A. Segerer, G. Tarmann, R. Zahiri, and H. C. J. Godfray

2016. Species-level para- and polyphyly in DNA barcode gene trees: Strong operational bias in european lepidoptera. *Systematic Biology*, 65(6):1024–1040.

812 Neal, R.
813 2011. MCMC using Hamiltonian dynamics. In *Handbook of Markov Chain Monte Carlo*.
814 Chapman & Hall, CRC Press.

815 Newcombe, R. G.
816 1998. Two-sided confidence intervals for the single proportion: comparison of seven
817 methods. *Statistics in Medicine*, 17(8):857–872.

818 Pante, E., N. Puillandre, A. Viricel, S. Arnaud-Haond, D. Aurelle, M. Castelin, A. Chemuil,
819 C. Destombe, D. Forcioli, M. Valero, F. Viard, and S. Samadi
820 2015. Species are hypotheses: Avoid connectivity assessments based on pillars of sand.
821 *Molecular Ecology*, 24(3):525–544.

822 Paradis, E., J. Claude, and K. Strimmer
823 2004. Ape: Analyses of phylogenetics and evolution in r language. *Bioinformatics*,
824 20(2):289–290.

825 Pentinsaari, M., H. Salmela, M. Mutanen, and T. Roslin
826 2016. Molecular evolution of a widely-adopted taxonomic marker (COI) across the animal
827 tree of life. *Scientific Reports*, 6:35275.

828 Phillips, J., T. Athey, P. McNicholas, and R. Hanner
829 2023. VLF: An r package for the analysis of very low frequency variants in DNA sequences.
830 *Biodiversity Data Journal*, 11:e96480.

831 Phillips, J., D. Gillis, and R. Hanner
832 2019. Incomplete estimates of genetic diversity within species: Implications for DNA
833 barcoding. *Ecology and Evolution*, 9(5):2996–3010.

834 Phillips, J., D. Gillis, and R. Hanner
835 2020. HACSIm: An R package to estimate intraspecific sample sizes for genetic diversity
836 assessment using haplotype accumulation curves. *PeerJ Computer Science*.

Phillips, J., D. Gillis, and R. Hanner

2022. Lack of statistical rigor in DNA barcoding likely invalidates the presence of a true species' barcode gap. *Frontiers in Ecology and Evolution*, 10:859099.

Phillips, J., C. Griswold, R. Young, N. Hubert, and H. Hanner

2024. *A Measure of the DNA Barcode Gap for Applied and Basic Research*, Pp. 375–390. New York, NY: Springer US.

Phillips, J., R. Gwiazdowski, D. Ashlock, and R. Hanner

2015. An exploration of sufficient sampling effort to describe intraspecific DNA barcode haplotype diversity: examples from the ray-finned fishes (Chordata: Actinopterygii). *DNA Barcodes*, 3:66–73.

Pons, J., T. G. Barraclough, J. Gomez-Zurita, A. Cardoso, D. P. Duran, S. Hazell, S. Kamoun, W. D. Sumlin, and A. P. Vogler

2006. Sequence-based species delimitation for the DNA taxonomy of undescribed insects. *Systematic Biology*, 55(4):595–609.

Puillandre, N., S. Brouillet, and G. Achaz

2021. ASAP: Assemble Species by Automatic Partitioning. *Molecular Ecology Resources*, 21(2):609–620.

Puillandre, N., A. Lambert, S. Brouillet, and G. Achaz

2011. ABGD, Automatic Barcode Gap Discovery for primary species delimitation. *Molecular Ecology*, 21(8):1864–1877.

R Core Team

2024. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria.

Rannala, B.

2015. The art and science of species delimitation. *Current Zoology*, 61(5):846–853.

Rannala, B. and Z. Yang

2003. Bayes estimation of species divergence times and ancestral population sizes using DNA sequences from multiple loci. *Genetics*, 164:1645–1656.

Rannala, B. and Z. Yang

2017. Efficient Bayesian species tree inference under the multispecies coalescent. *Systematic Biology*, 66(5):823–842.

Ratnasingham, S. and P. Hebert

2007. BOLD: The Barcode of Life Data System (<http://www.barcodinglife.org>). *Molecular Ecology Notes*, 7(3):355–364.

Ratnasingham, S. and P. D. N. Hebert

2013. A DNA-based registry for all animal species: The barcode index number (BIN) system. *PLoS One*, 8(7):e66213.

Rosenberg, N. and M. Nordborg

2002. Genealogical trees, coalescent theory and the analysis of genetic polymorphisms. *Nature Reviews Genetics*, 3(5):380–390.

Stan Development Team

2023. RStan: The R interface to Stan. R package version 2.32.6.

Stephens, M.

2000. Dealing with label switching in mixture models. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 62(4):795–809.

Stoeckle, M. and D. Thaler

2014. DNA barcoding works in practice but not in (neutral) theory. *PLoS One*, 9(7):e100755.

885 Talavera, G., V. Dincă, and R. Vila

886 2013. Factors affecting species delimitations with the GMYC model: insights from a
887 butterfly survey. *Methods in Ecology and Evolution*, 4(12):1101–1110.

888 Templeton, A., K. Crandall, and C. Sing

889 1992. A cladistic analysis of phenotypic associations with haplotypes inferred from
890 restriction endonuclease mapping and DNA sequence data. III. Cladogram estimation.
891 *Genetics*, 132(2):619–633.

892 Vehtari, A., A. Gelman, D. Simpson, B. Carpenter, and P.-C. Bürkner

893 2021. Rank-normalization, folding, and localization: An improved $\hat{R}$ for assessing
894 convergence of MCMC (with discussion). *Bayesian Analysis*, 16(2):667–718.

895 Wells, T., T. Carruthers, P. Muñoz-Rodríguez, A. Sumadijaya, J. R. I. Wood, and R. W.
896 Scotland

897 2022. Species as a heuristic: Reconciling theory and practice. *Systematic Biology*,
898 71(5):1233–1243.

899 Wickham, H.

900 2016. *ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York.

901 Wiemers, M. and K. Fiedler

902 2007. Does the DNA barcoding gap exist? - a case study in blue butterflies (Lepidoptera:
903 Lycaenidae). *Frontiers in Zoology*, 4(8).

904 Yang, Z. and B. Rannala

905 2010. Bayesian species delimitation using multilocus sequence data. *Proceedings of the*
906 *National Academy of Sciences*, 107:9264–9269.

907 Yang, Z. and B. Rannala

908 2014. Unguided species delimitation using dna sequence data from multiple loci. *Molecular*
909 *Biology and Evolution*, 31(12):3125–3135.

910 Yang, Z. and B. Rannala
911 2017. Bayesian species identification under the multispecies coalescent provides significant
912 improvements to DNA barcoding analyses. *Molecular Ecology*, 26:3028–3036.

913 Young, R., R. Gill, D. Gillis, and R. Hanner
914 2021. Molecular Acquisition, Cleaning and Evaluation in R (MACER) - A tool to assemble
915 molecular marker datasets from BOLD and GenBank. *Biodiversity Data Journal*, 9:e71378.

916 Zhang, J., P. Kapli, P. Pavlidis, and A. Stamatakis
917 2013. A general species delimitation method with applications to phylogenetic placements.
918 *Bioinformatics*, 29(22):2869–2876.